

GDC®

Circular Separable Convolution Depth of Field “Circular Dof”

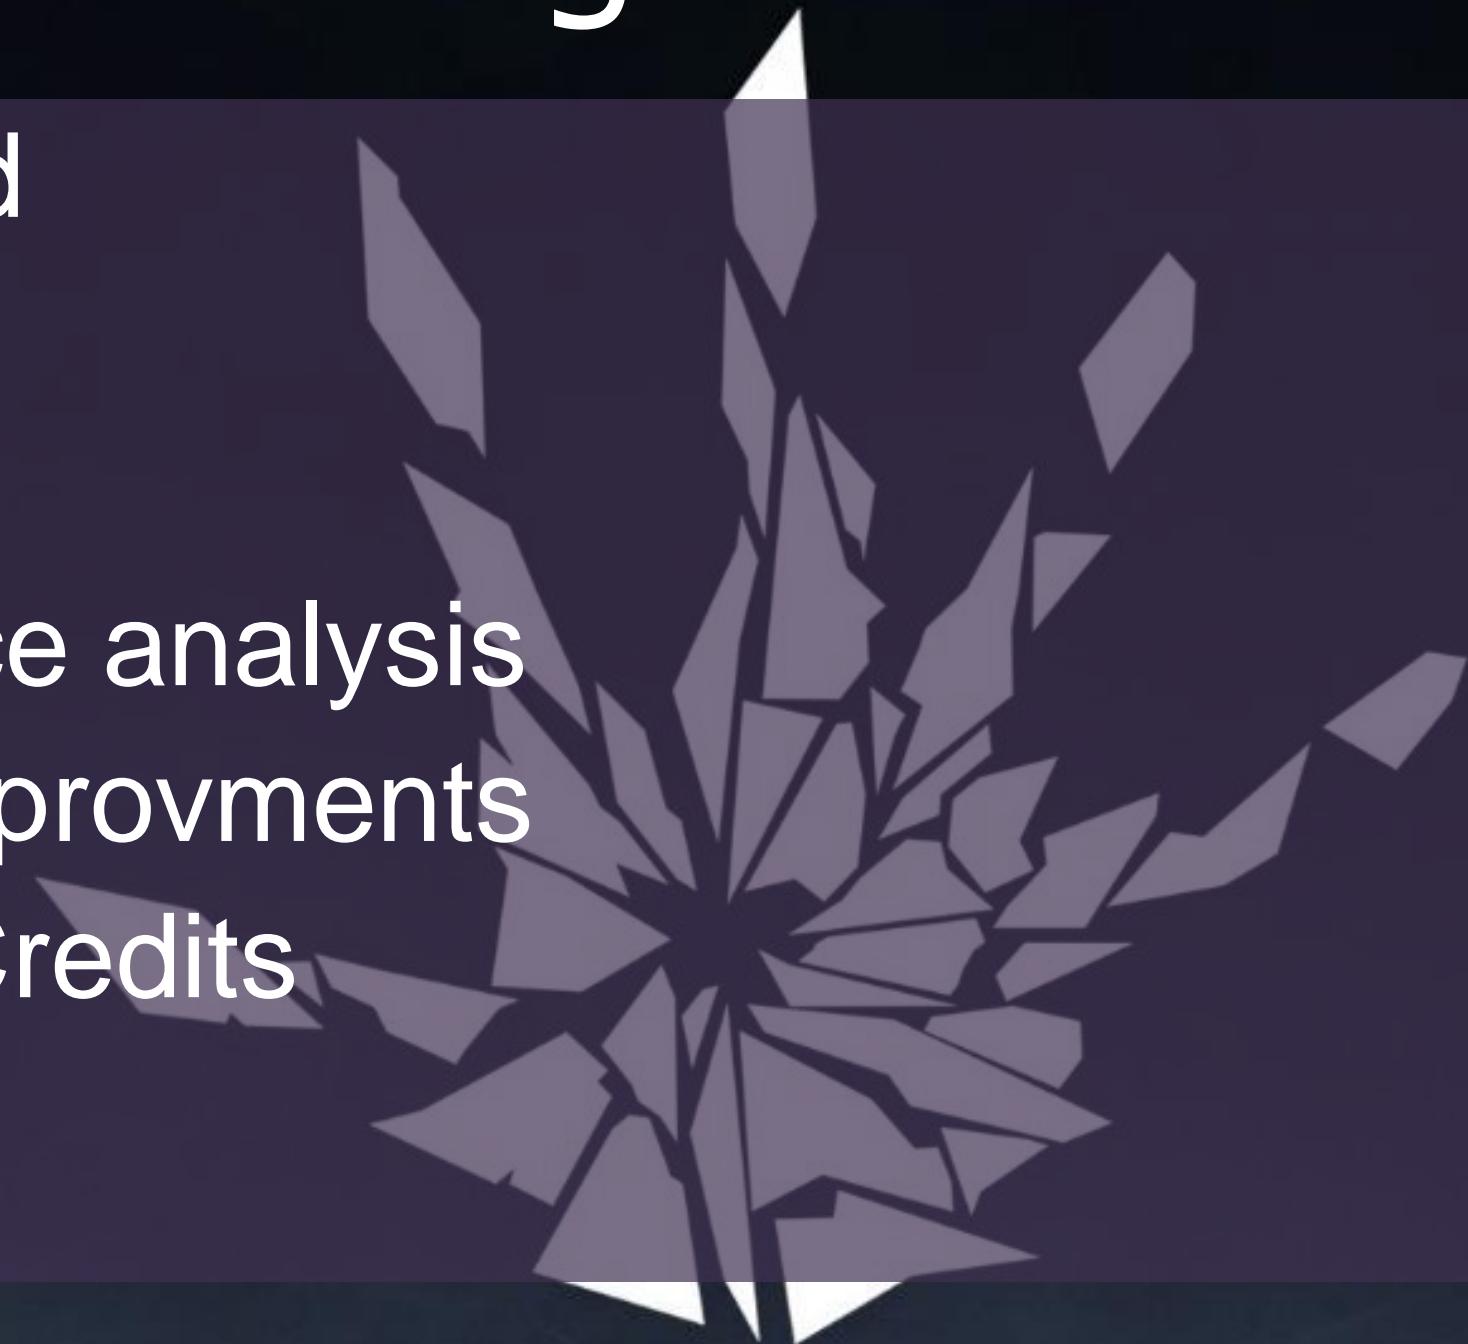
Kleber Garcia
Rendering Engineer – Frostbite Engine
Electronic Arts Inc.

GAME DEVELOPERS CONFERENCE® | MARCH 19-23, 2018 | EXPO: MARCH 21-23, 2018 #GDC18

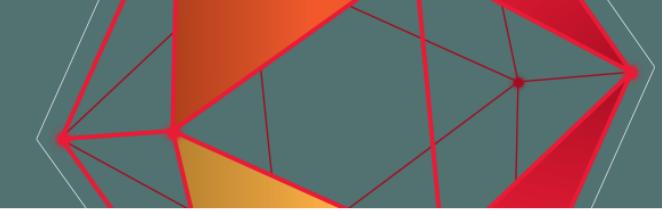


Agenda

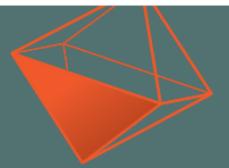
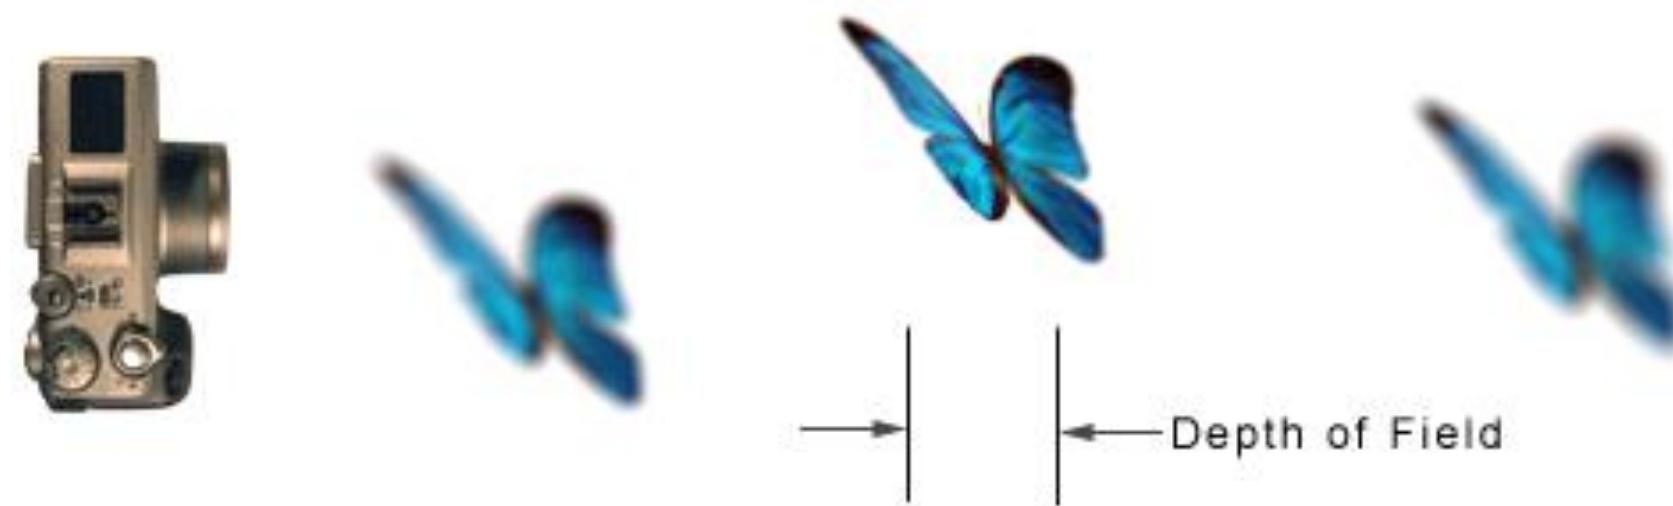
- Background
- Results
- Algorithm
- Performance analysis
- Artifacts/Improvements
- Sources / Credits
- Q/A

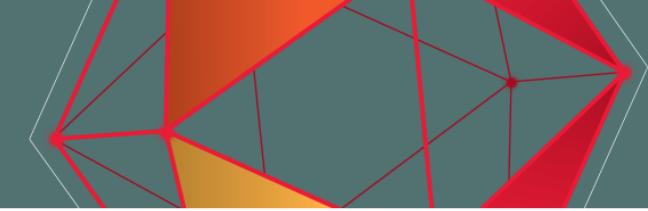


FROSTBITE™

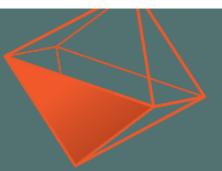
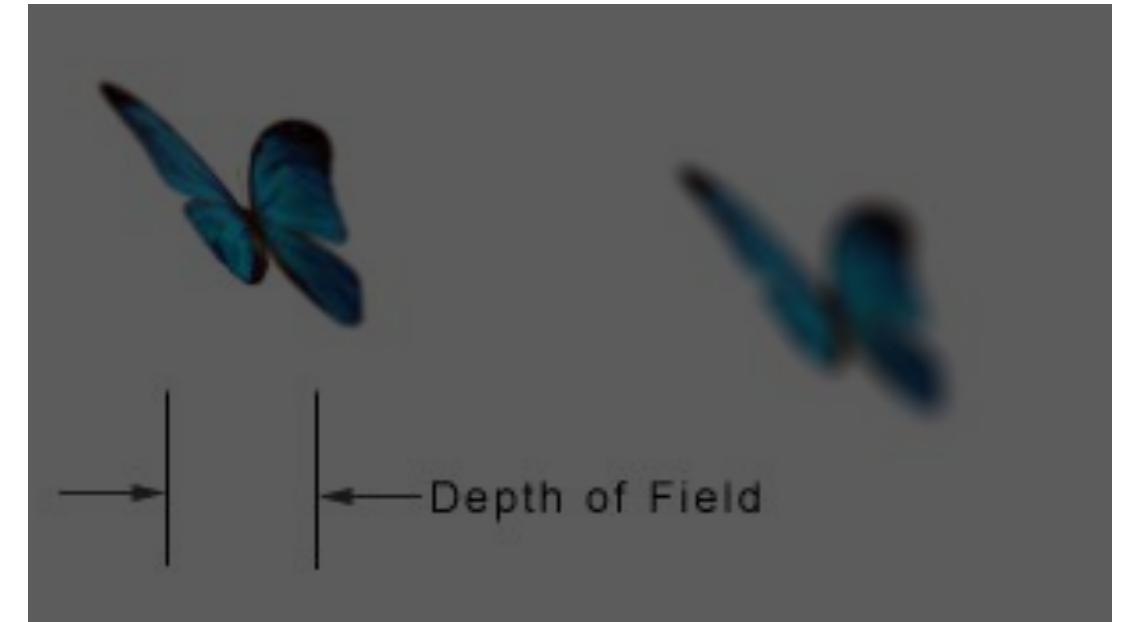


Background



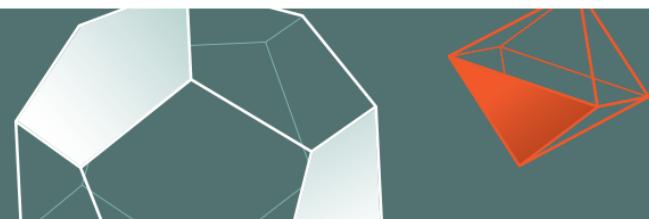
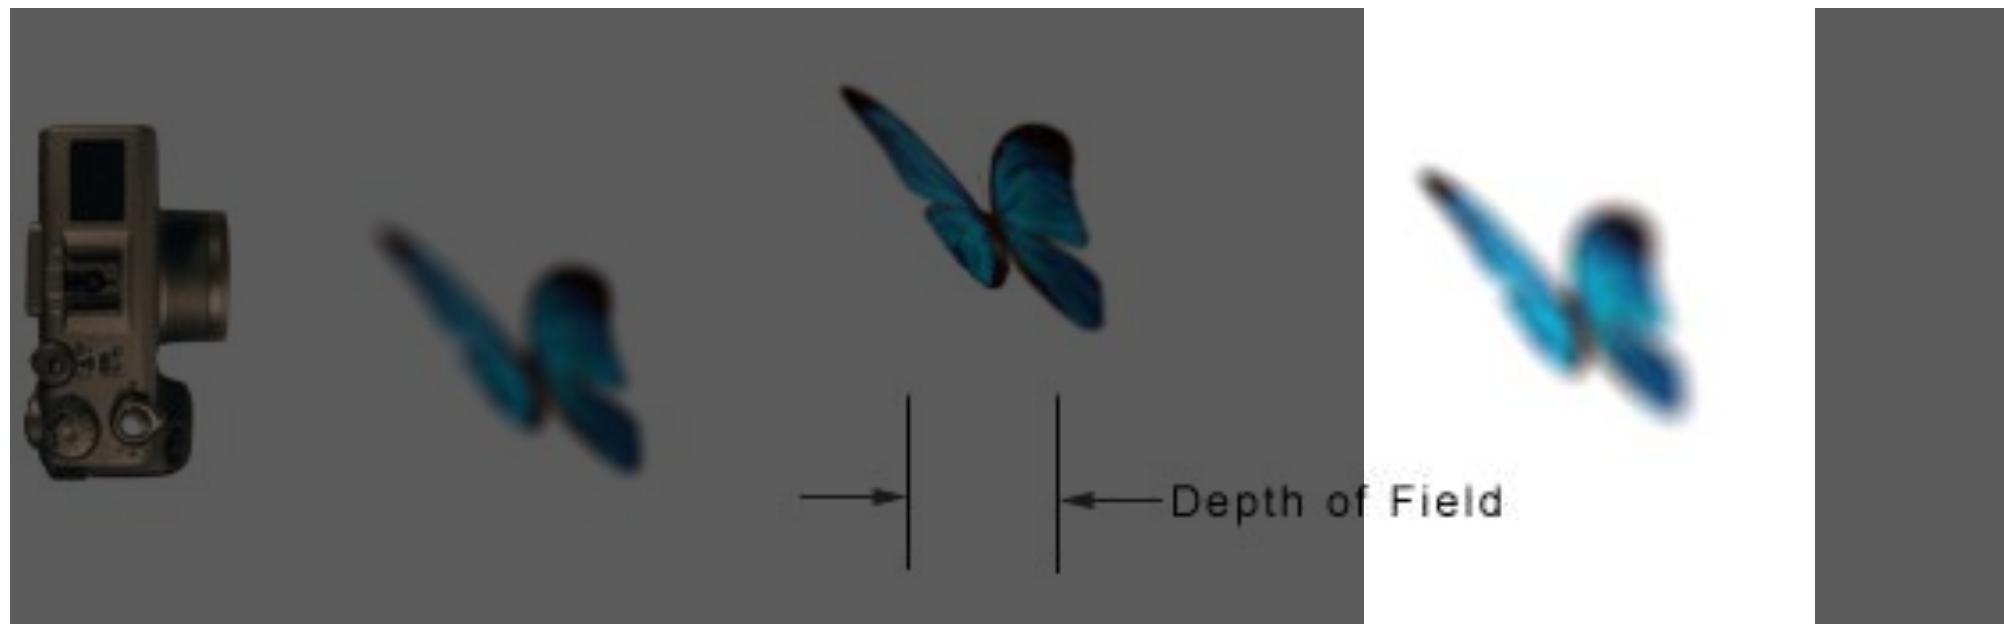


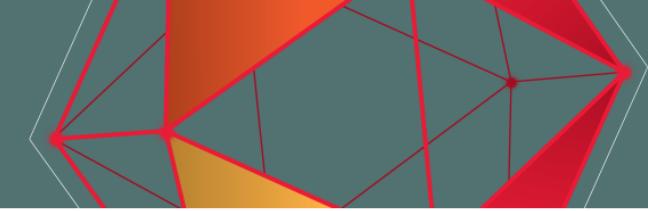
Background





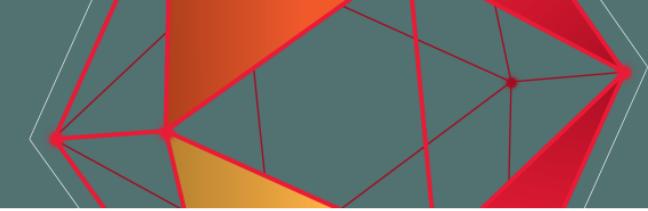
Background





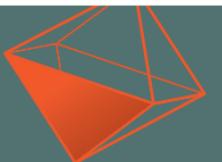
Background





Background

- Circle of Confusion (COC) – optical spot caused by a cone of light rays from a lens not coming to a perfect focus when imaging a point source.
- Can be thought of the 'radius of the blur' at a given pixel.



Visual Target



MADDEN





EA
SPORTS

MADDEN









FROSTBITE



FIFA 17



Sprite DOF

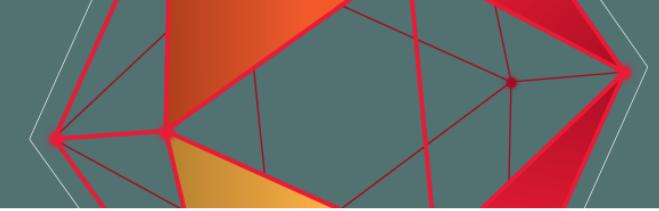
A close-up photograph of a dark, textured wooden surface. In the foreground, a small, dark mug with a decorative pattern sits next to a metal chain with large, spherical links. The background is blurred, showing more of the wooden structure and some bright, out-of-focus lights.

Circular DOF

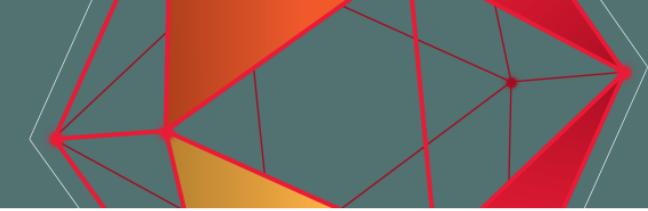
GDC®

GAME DEVELOPERS CONFERENCE®

| MARCH 19-23, 2018 | EXPO: MARCH 21-23, 2018 #GDC18



UBM



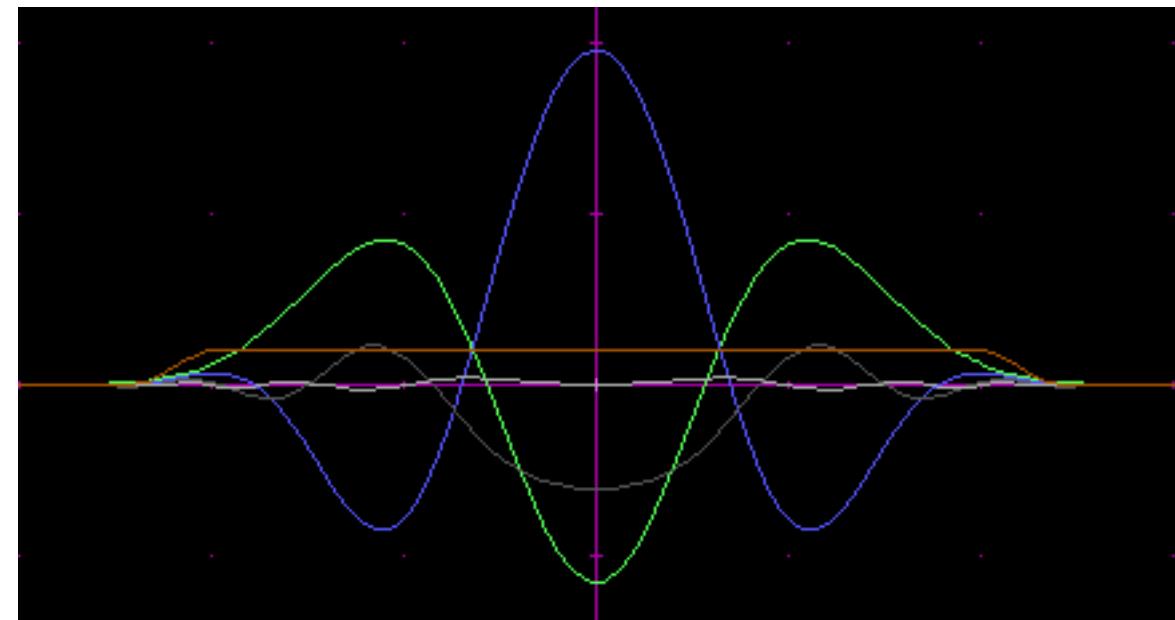
Algorithm

Separable gather (Circular Dof filter)

- <http://yehar.com/blog/?p=1495>

By Olli Niemitalo

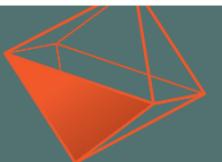
- Separable circular filter.
- Possible in the frequency domain (imaginary space)!
- Decompose frame buffer into a Fourier Transform
 - Possible to mix the signals and get a circular convolution





Algorithm

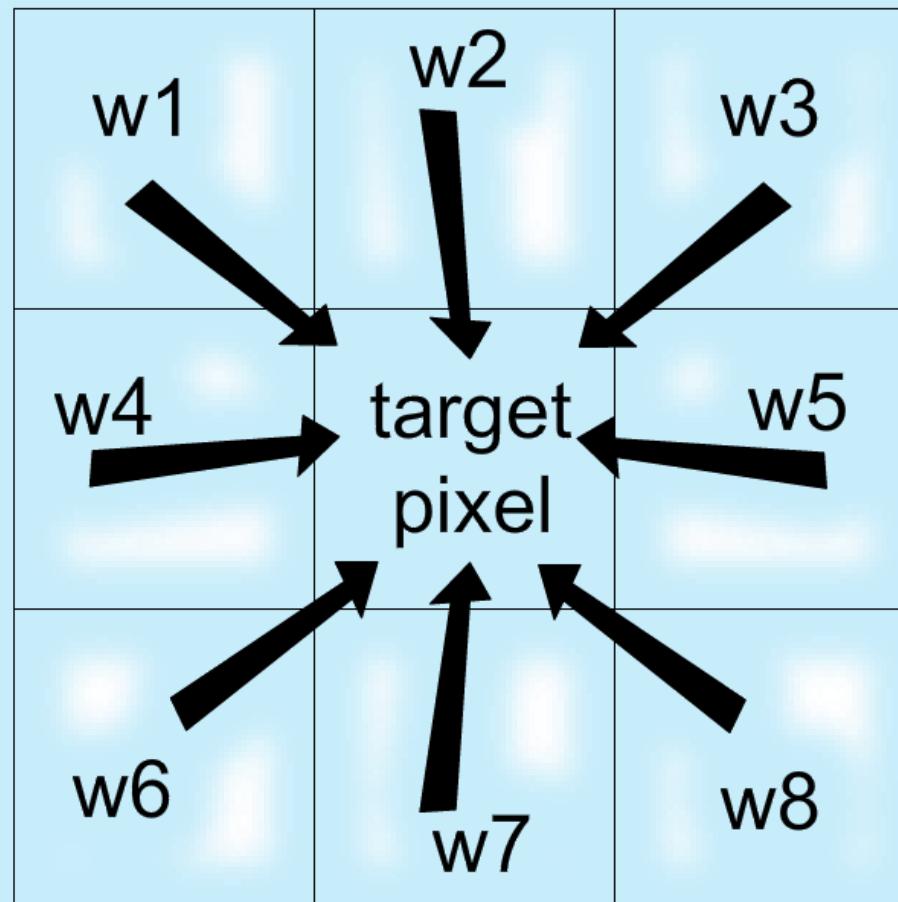
- To understand the separable property of Circular Dof, lets first take a look at how separable Gather works.



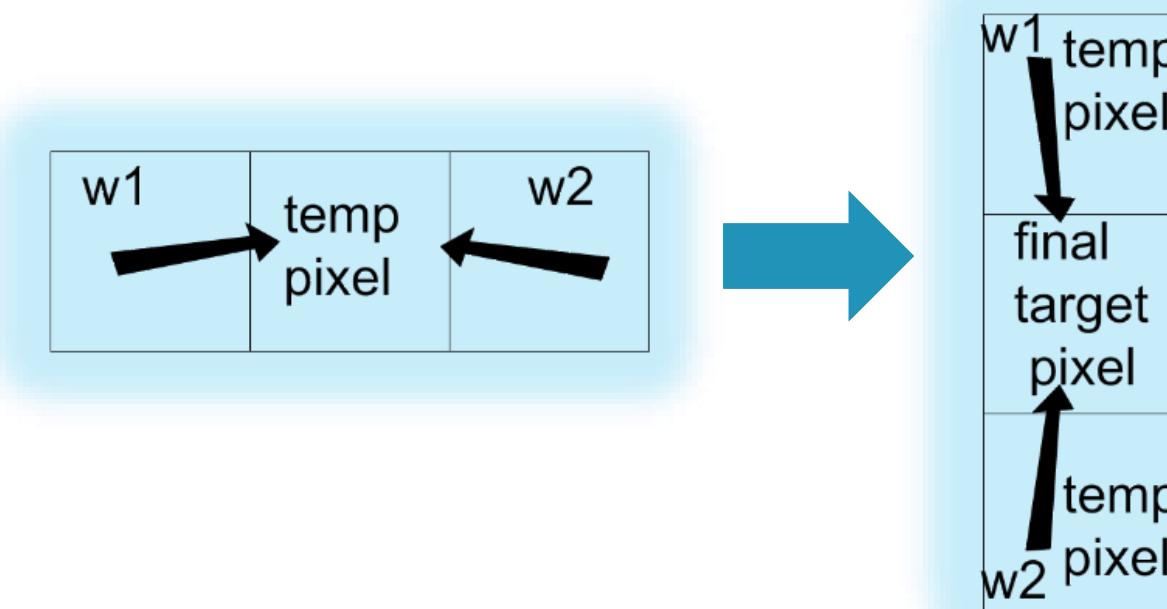


Algorithm

Brute force gather vs Separable Gather



Brute Force Gather - $O(n^2)$



Separable Gather - $O(n)$

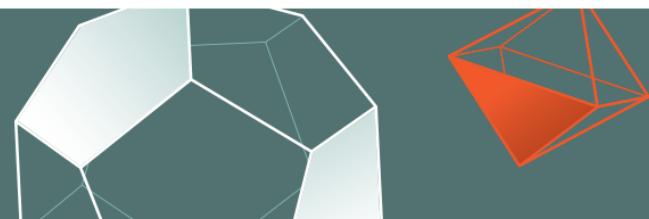


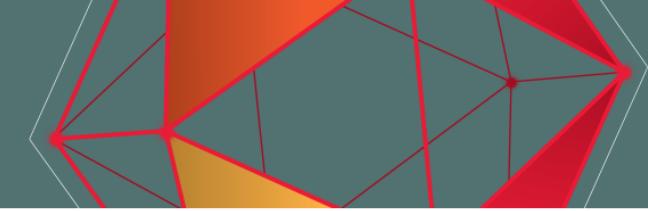


Algorithm

Separable Bokeh

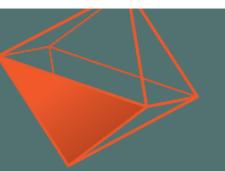
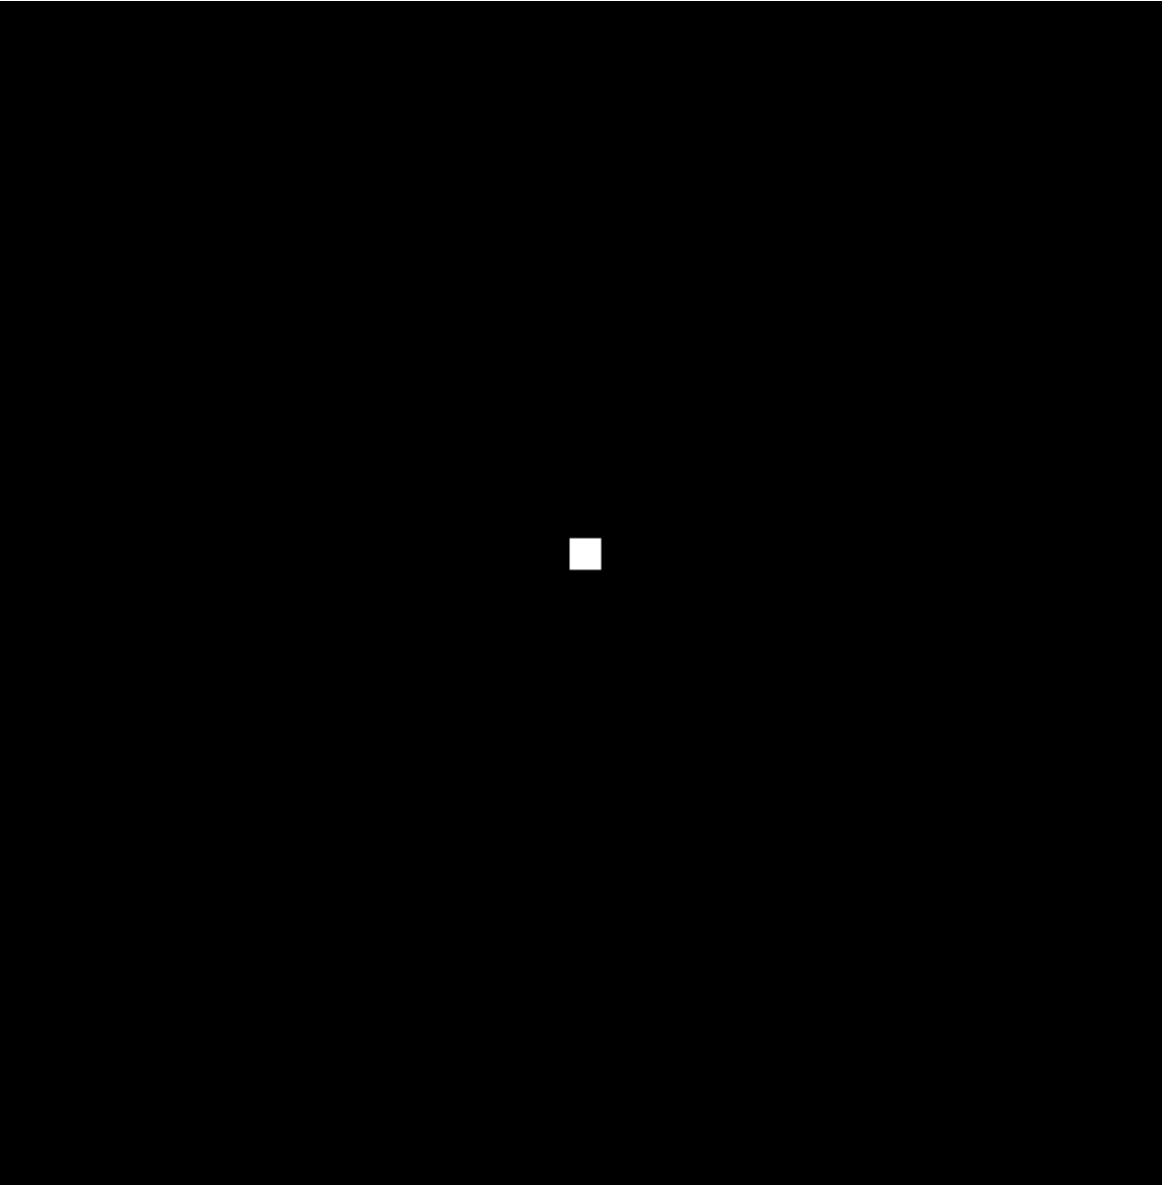
- Our approach has same time complexity as separable Gather-Gaussian.





Algorithm

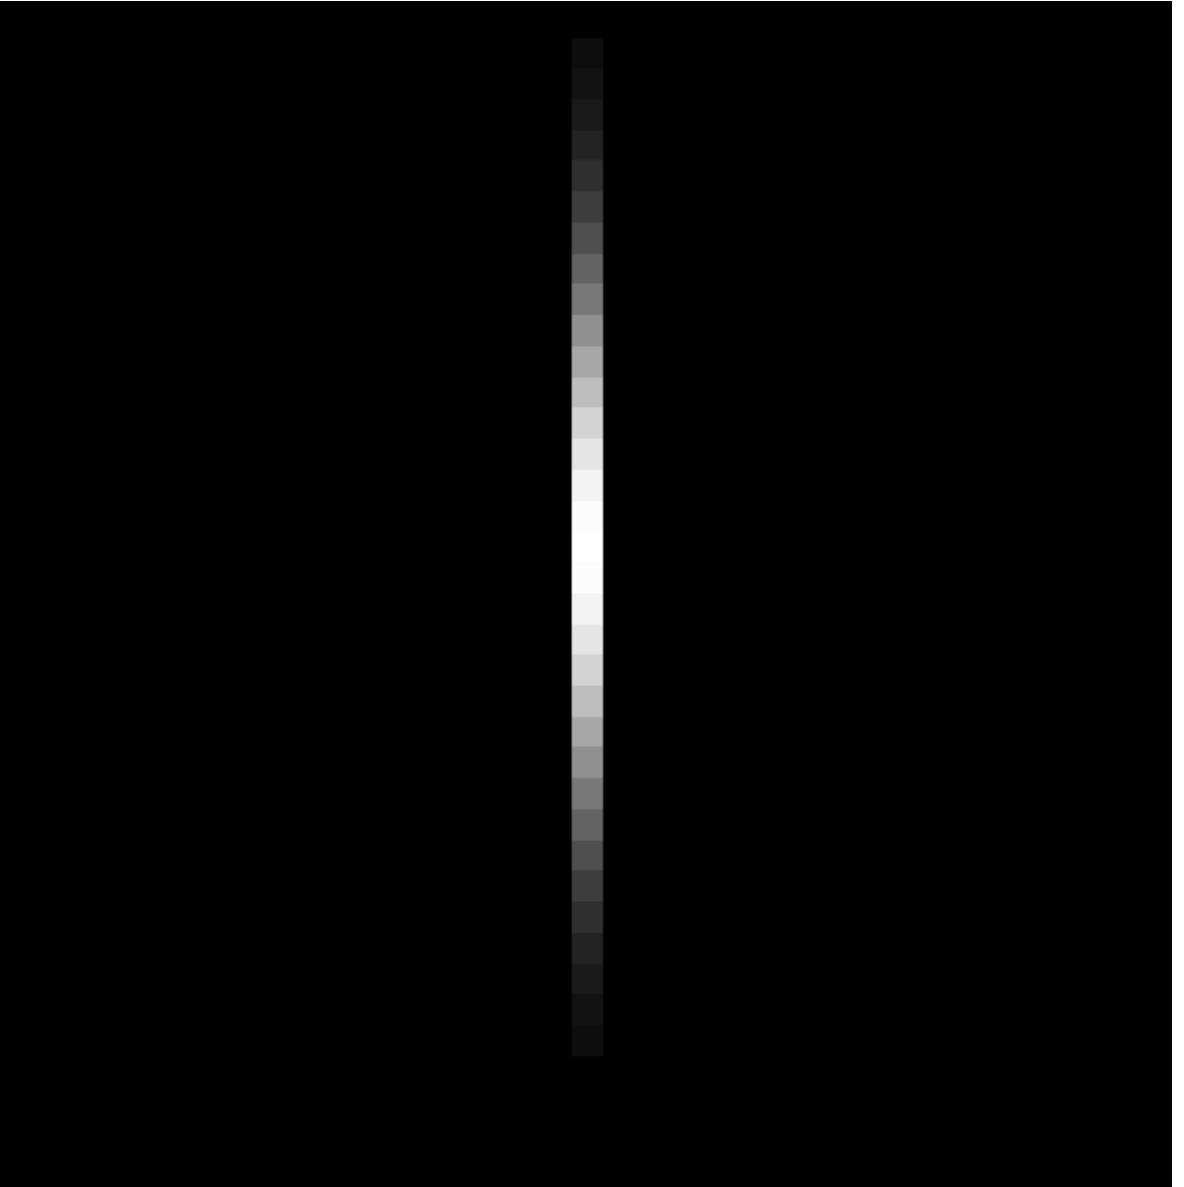
Separable gather (Gaussian filter)





Algorithm

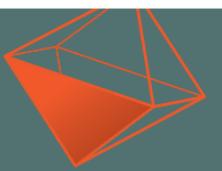
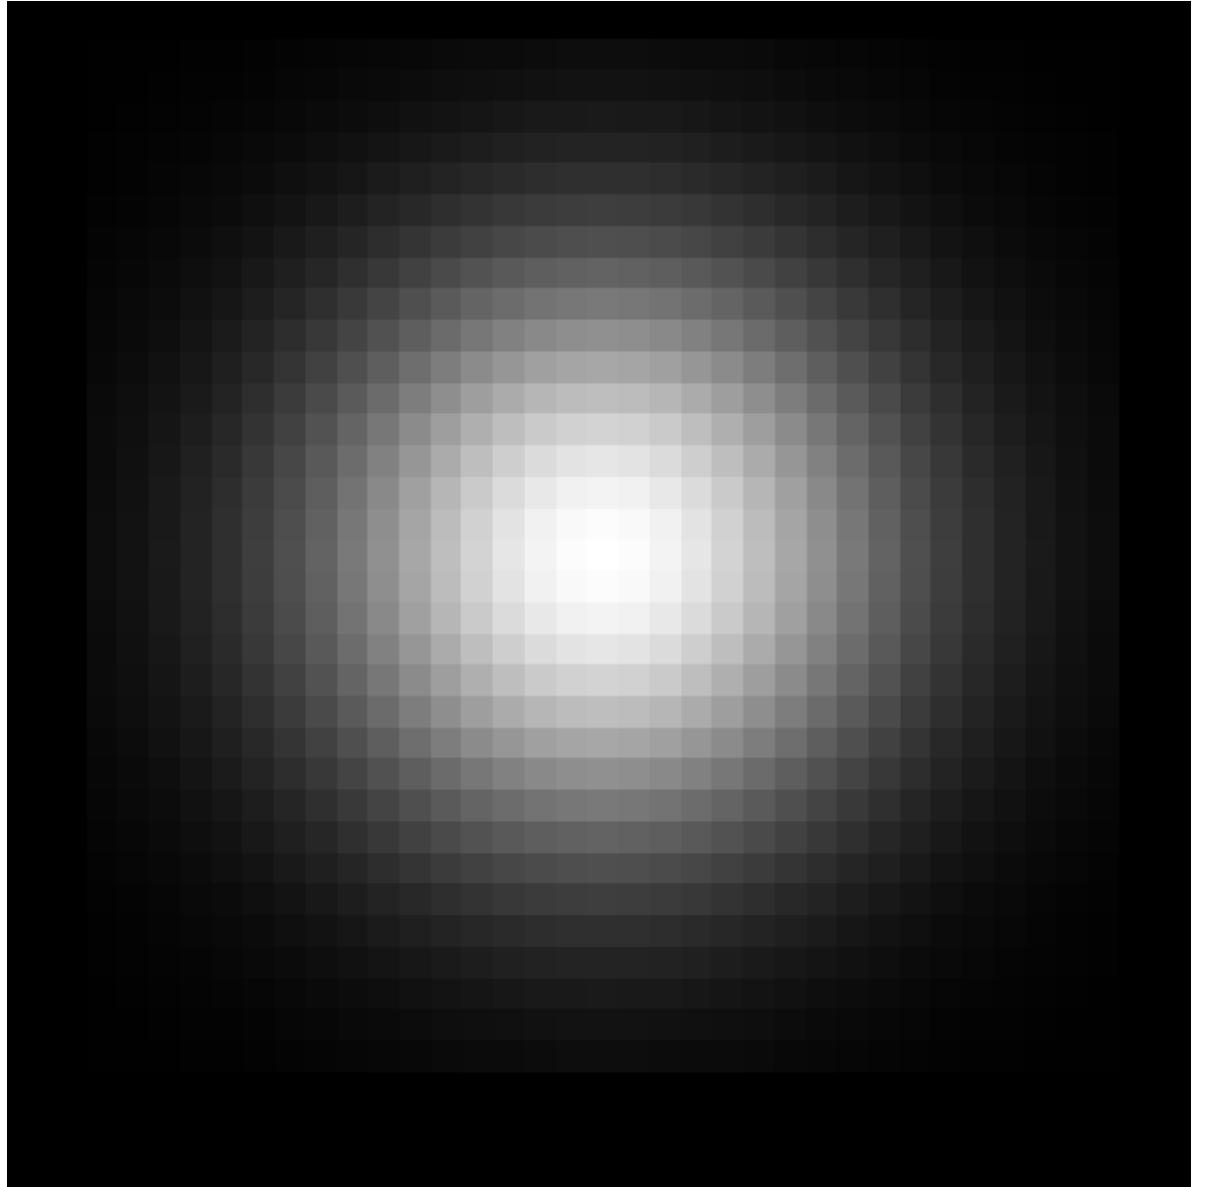
Separable gather (Gaussian filter)





Algorithm

Separable gather (Gaussian filter)





Algorithm

Separable gather (Gaussian filter)

- $F(x) = e^{-ax^2}$



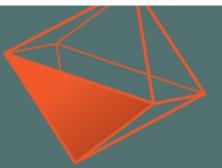
Clear Image



Vertical Blur



Horizontal Blur





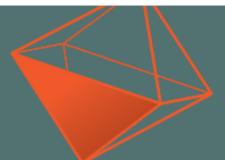
Algorithm

- A filter $F(x)$ can be resolved to a set of weights.
- Our approach resolves a complex filter into a complex number
- Complex numbers have 2 components, real and imaginary
- Remember $i * i = -1$
- Let P be a complex number, $P = P_r + P_i i$
- The sum of two complex numbers P and Q would be

$$P + Q = (P_r + Q_r) + (P_i + Q_i)i$$

- The product of two complex numbers would be

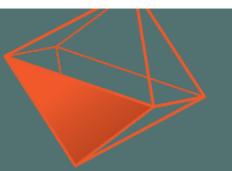
$$P * Q = (P_r * Q_r) - (P_i * Q_i) + [(P_r * Q_i) + (Q_r * P_i)]i$$





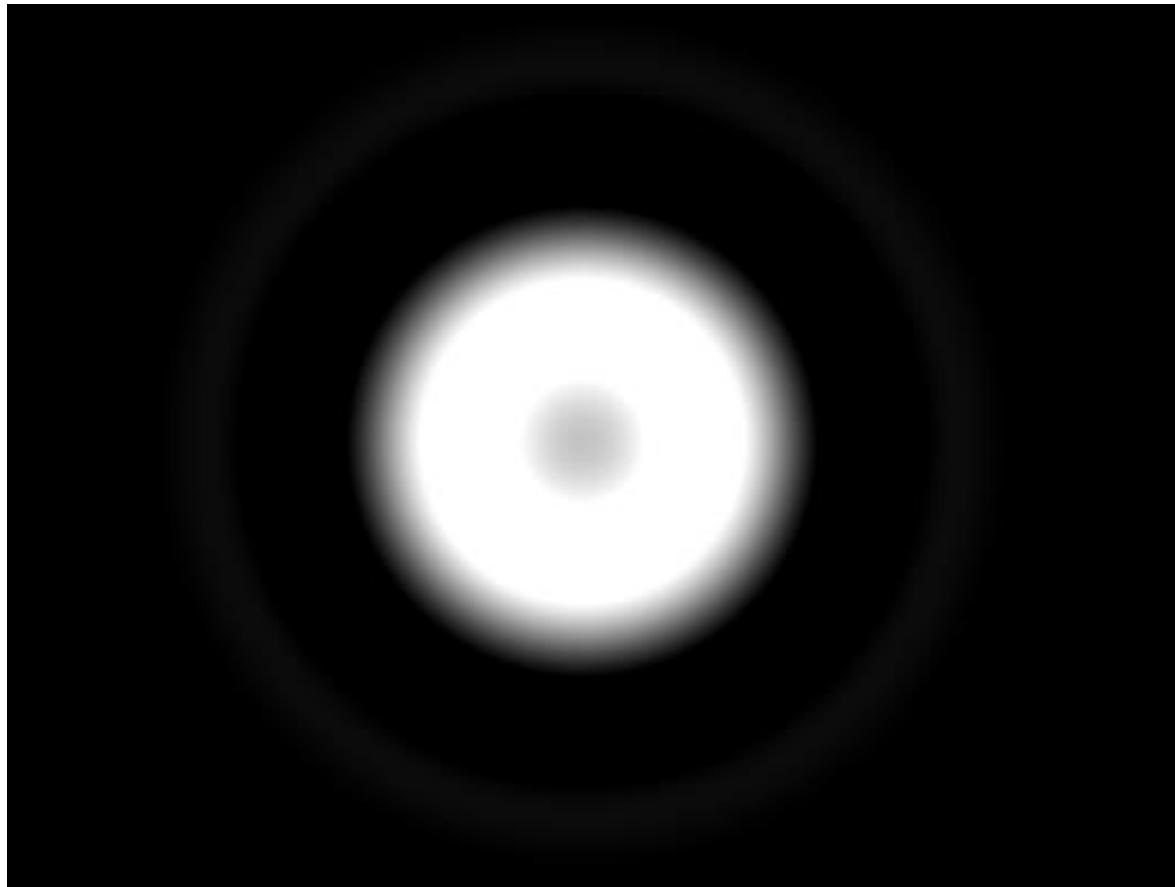
Algorithm

- Lets look now at circular DoF in action...

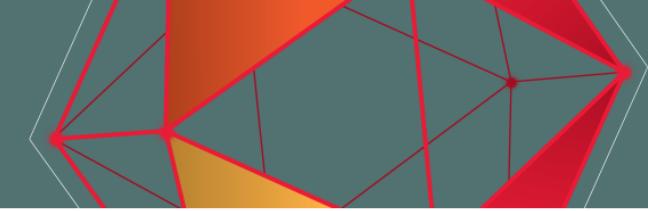


Algorithm

Separable gather (Circular Dof filter)

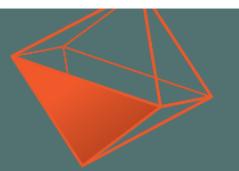
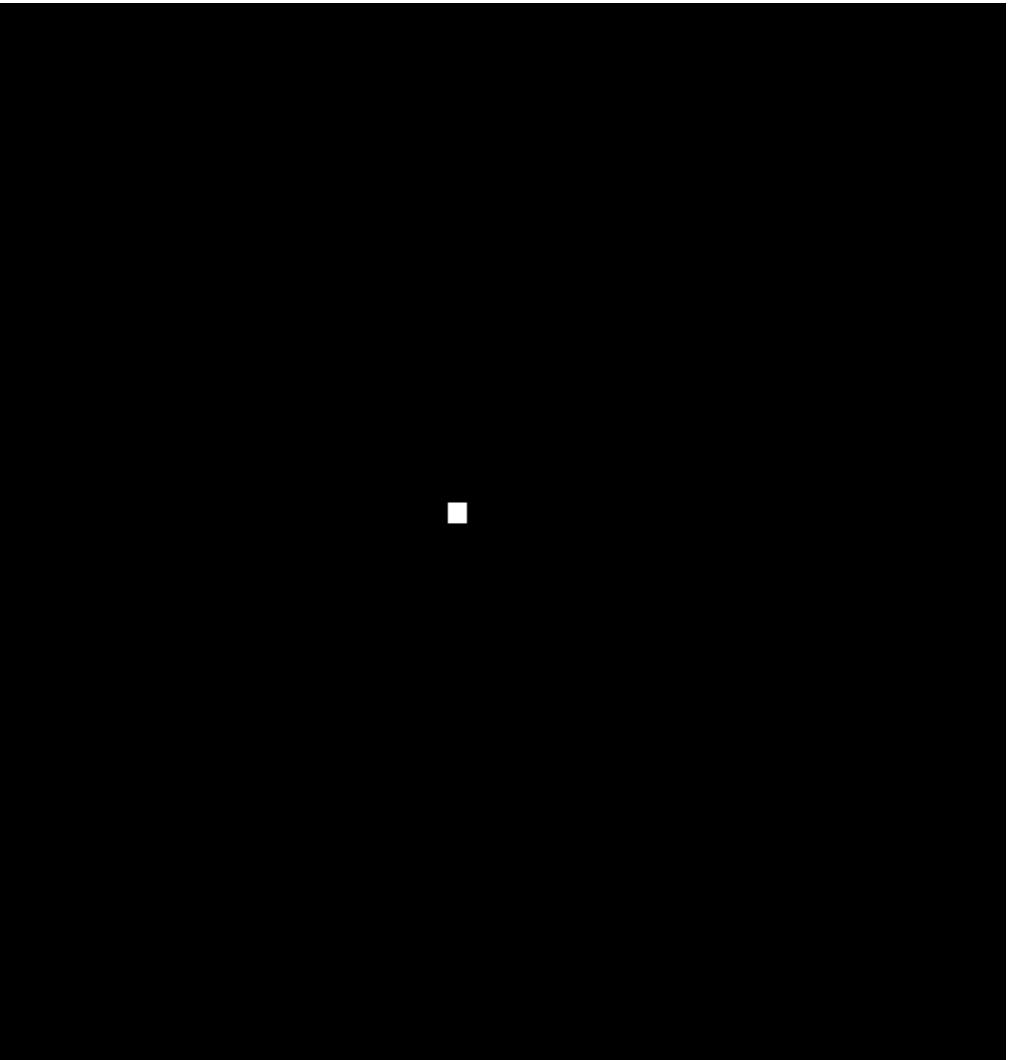


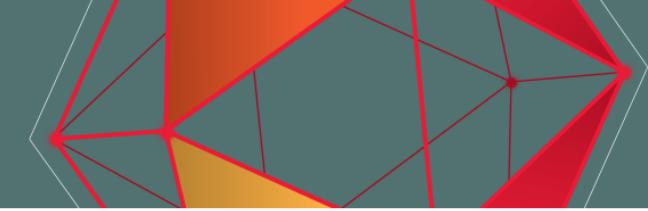
One Component Filter



Algorithm

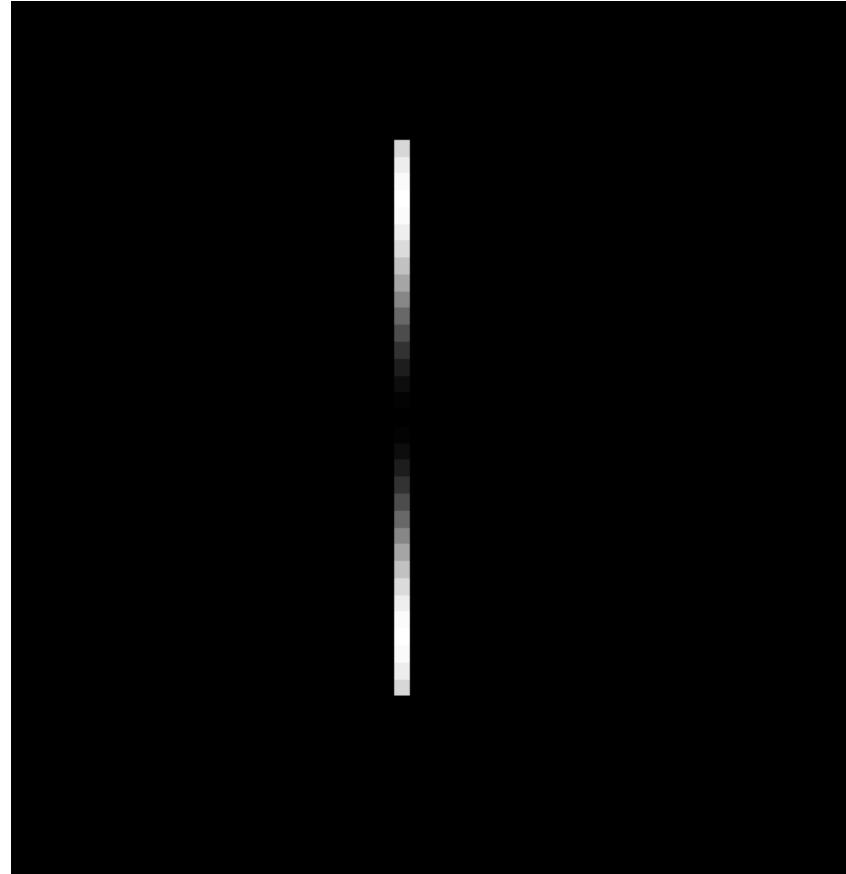
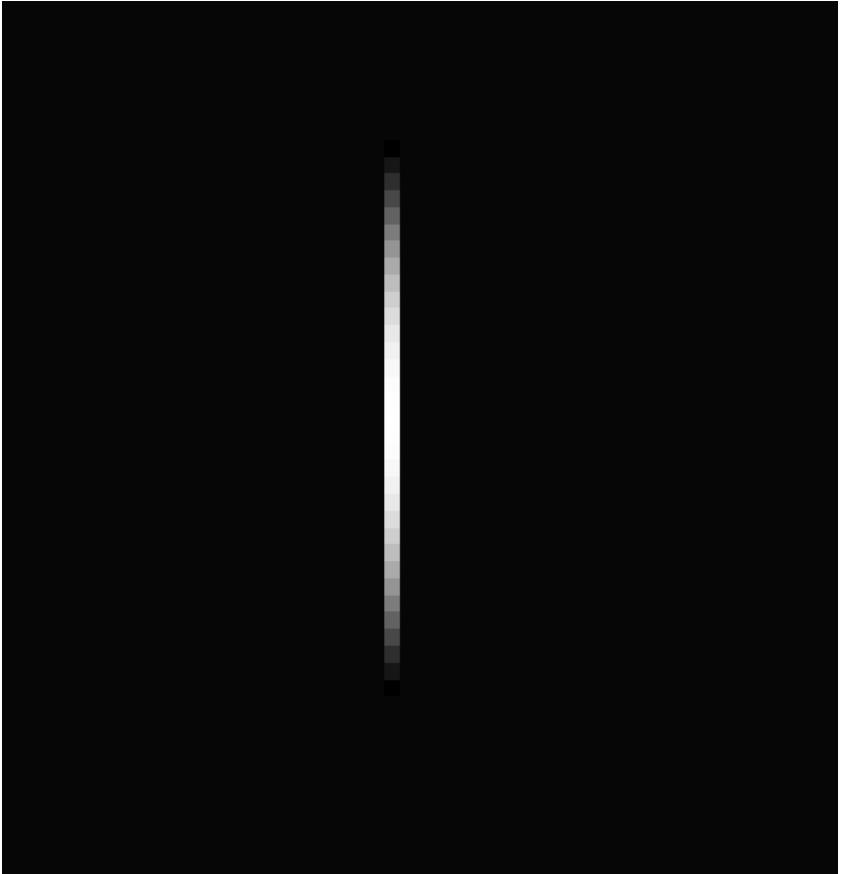
Separable gather (Circular Dof filter)



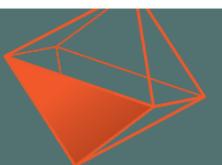


Algorithm

Separable gather (Circular Dof filter)



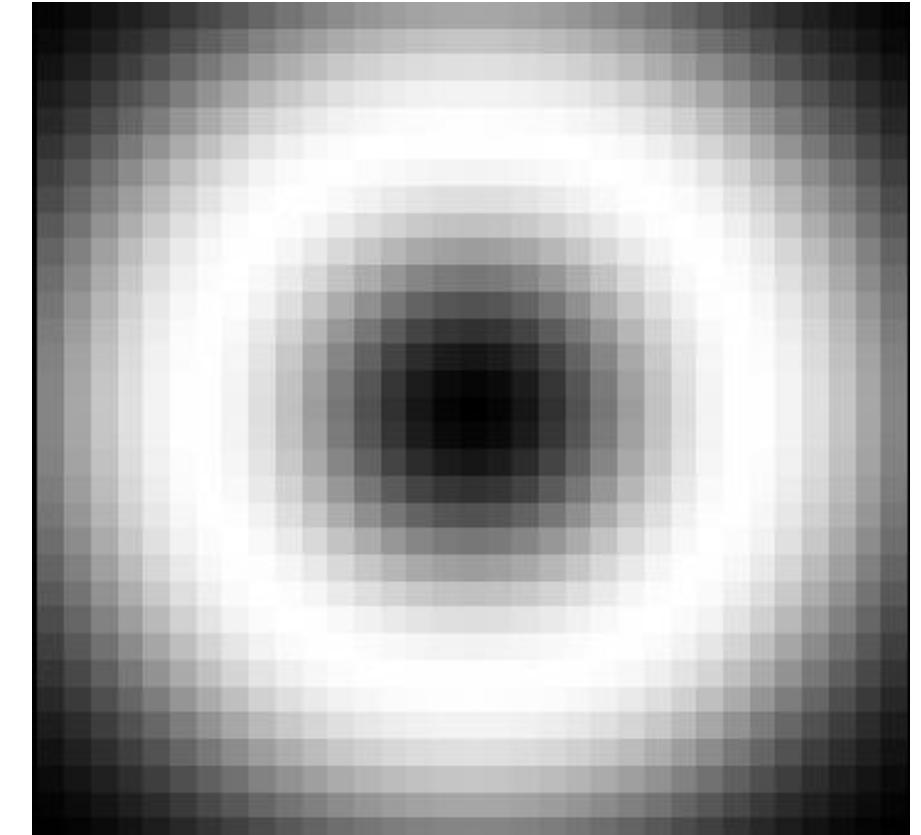
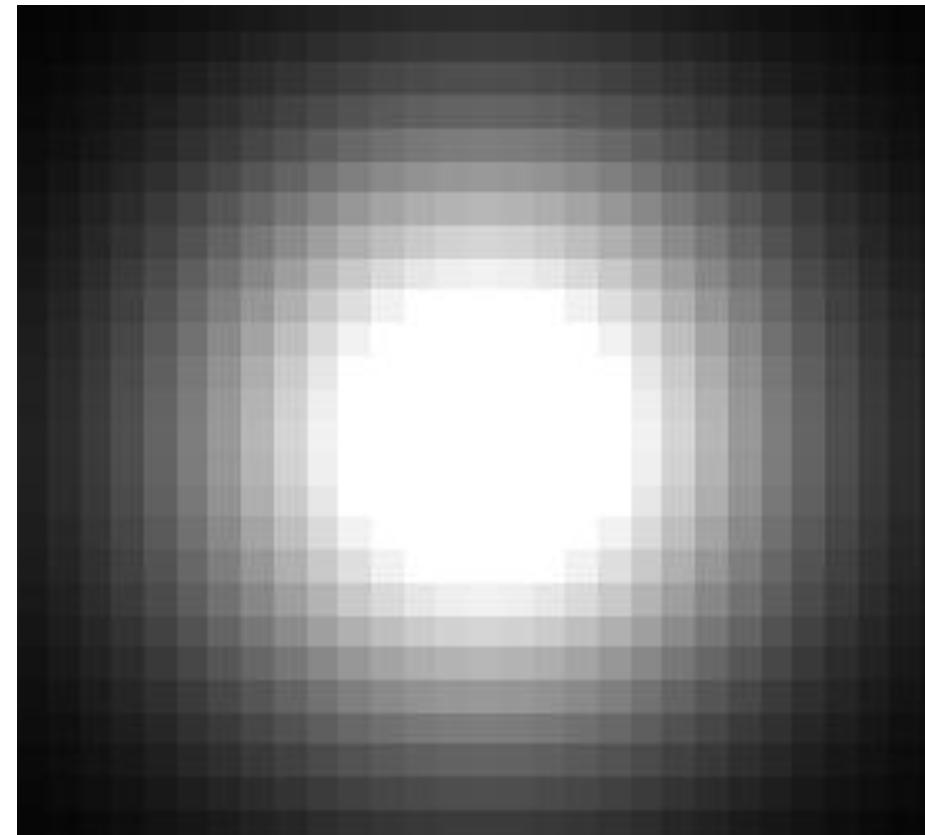
$$F(x) = e^{-ax^2}(\cos(bx^2) + i \sin(bx^2))$$



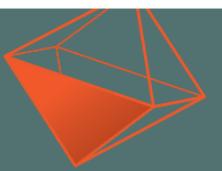


Algorithm

Separable gather (Circular Dof filter)



$$F(x) = e^{-ax^2}(\cos(bx^2) + i \sin(bx^2))$$





Algorithm

Separable gather (Circular Dof filter)



$$\text{Color}(x) = A * F_{real}(x) + B * F_{imaginary}(x)$$



Circular Dof: $F(x) = e^{-ax^2}(\cos(bx^2) + i \sin(bx^2))$

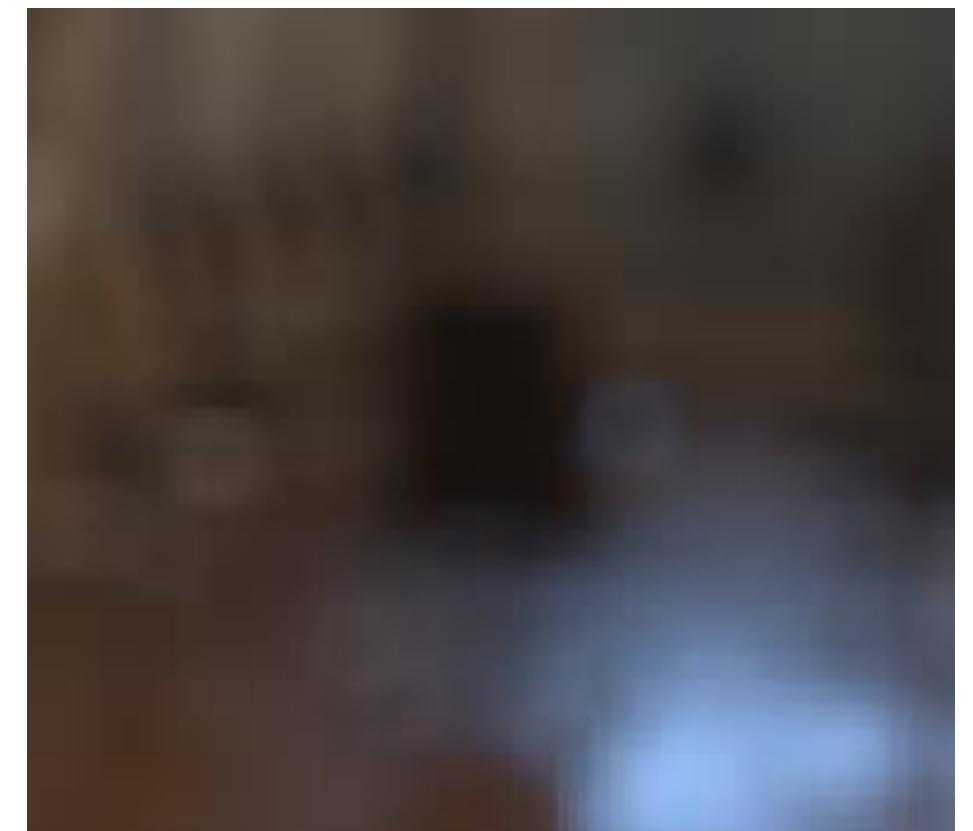


Real component

Vertical Blur



Imaginary component



Final Image

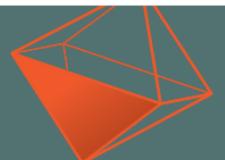
Horizontal blur & combine

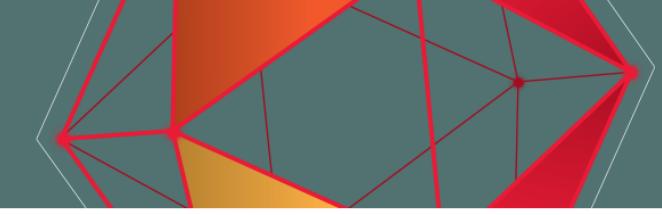


Algorithm

Separable gather (Circular Dof filter)

- That was just 1 component. We can add filters (multiple components) and approximate a circle better.

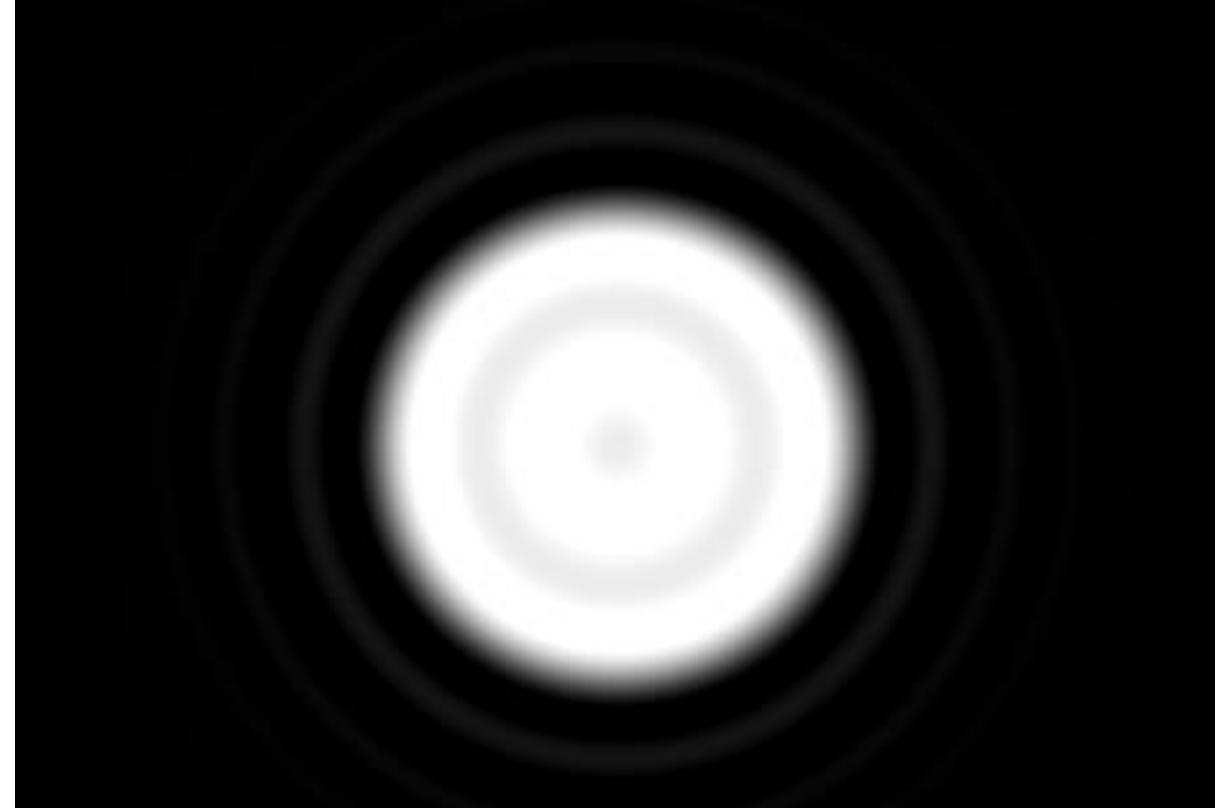




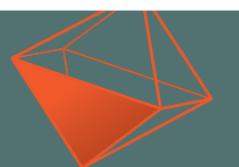
Algorithm

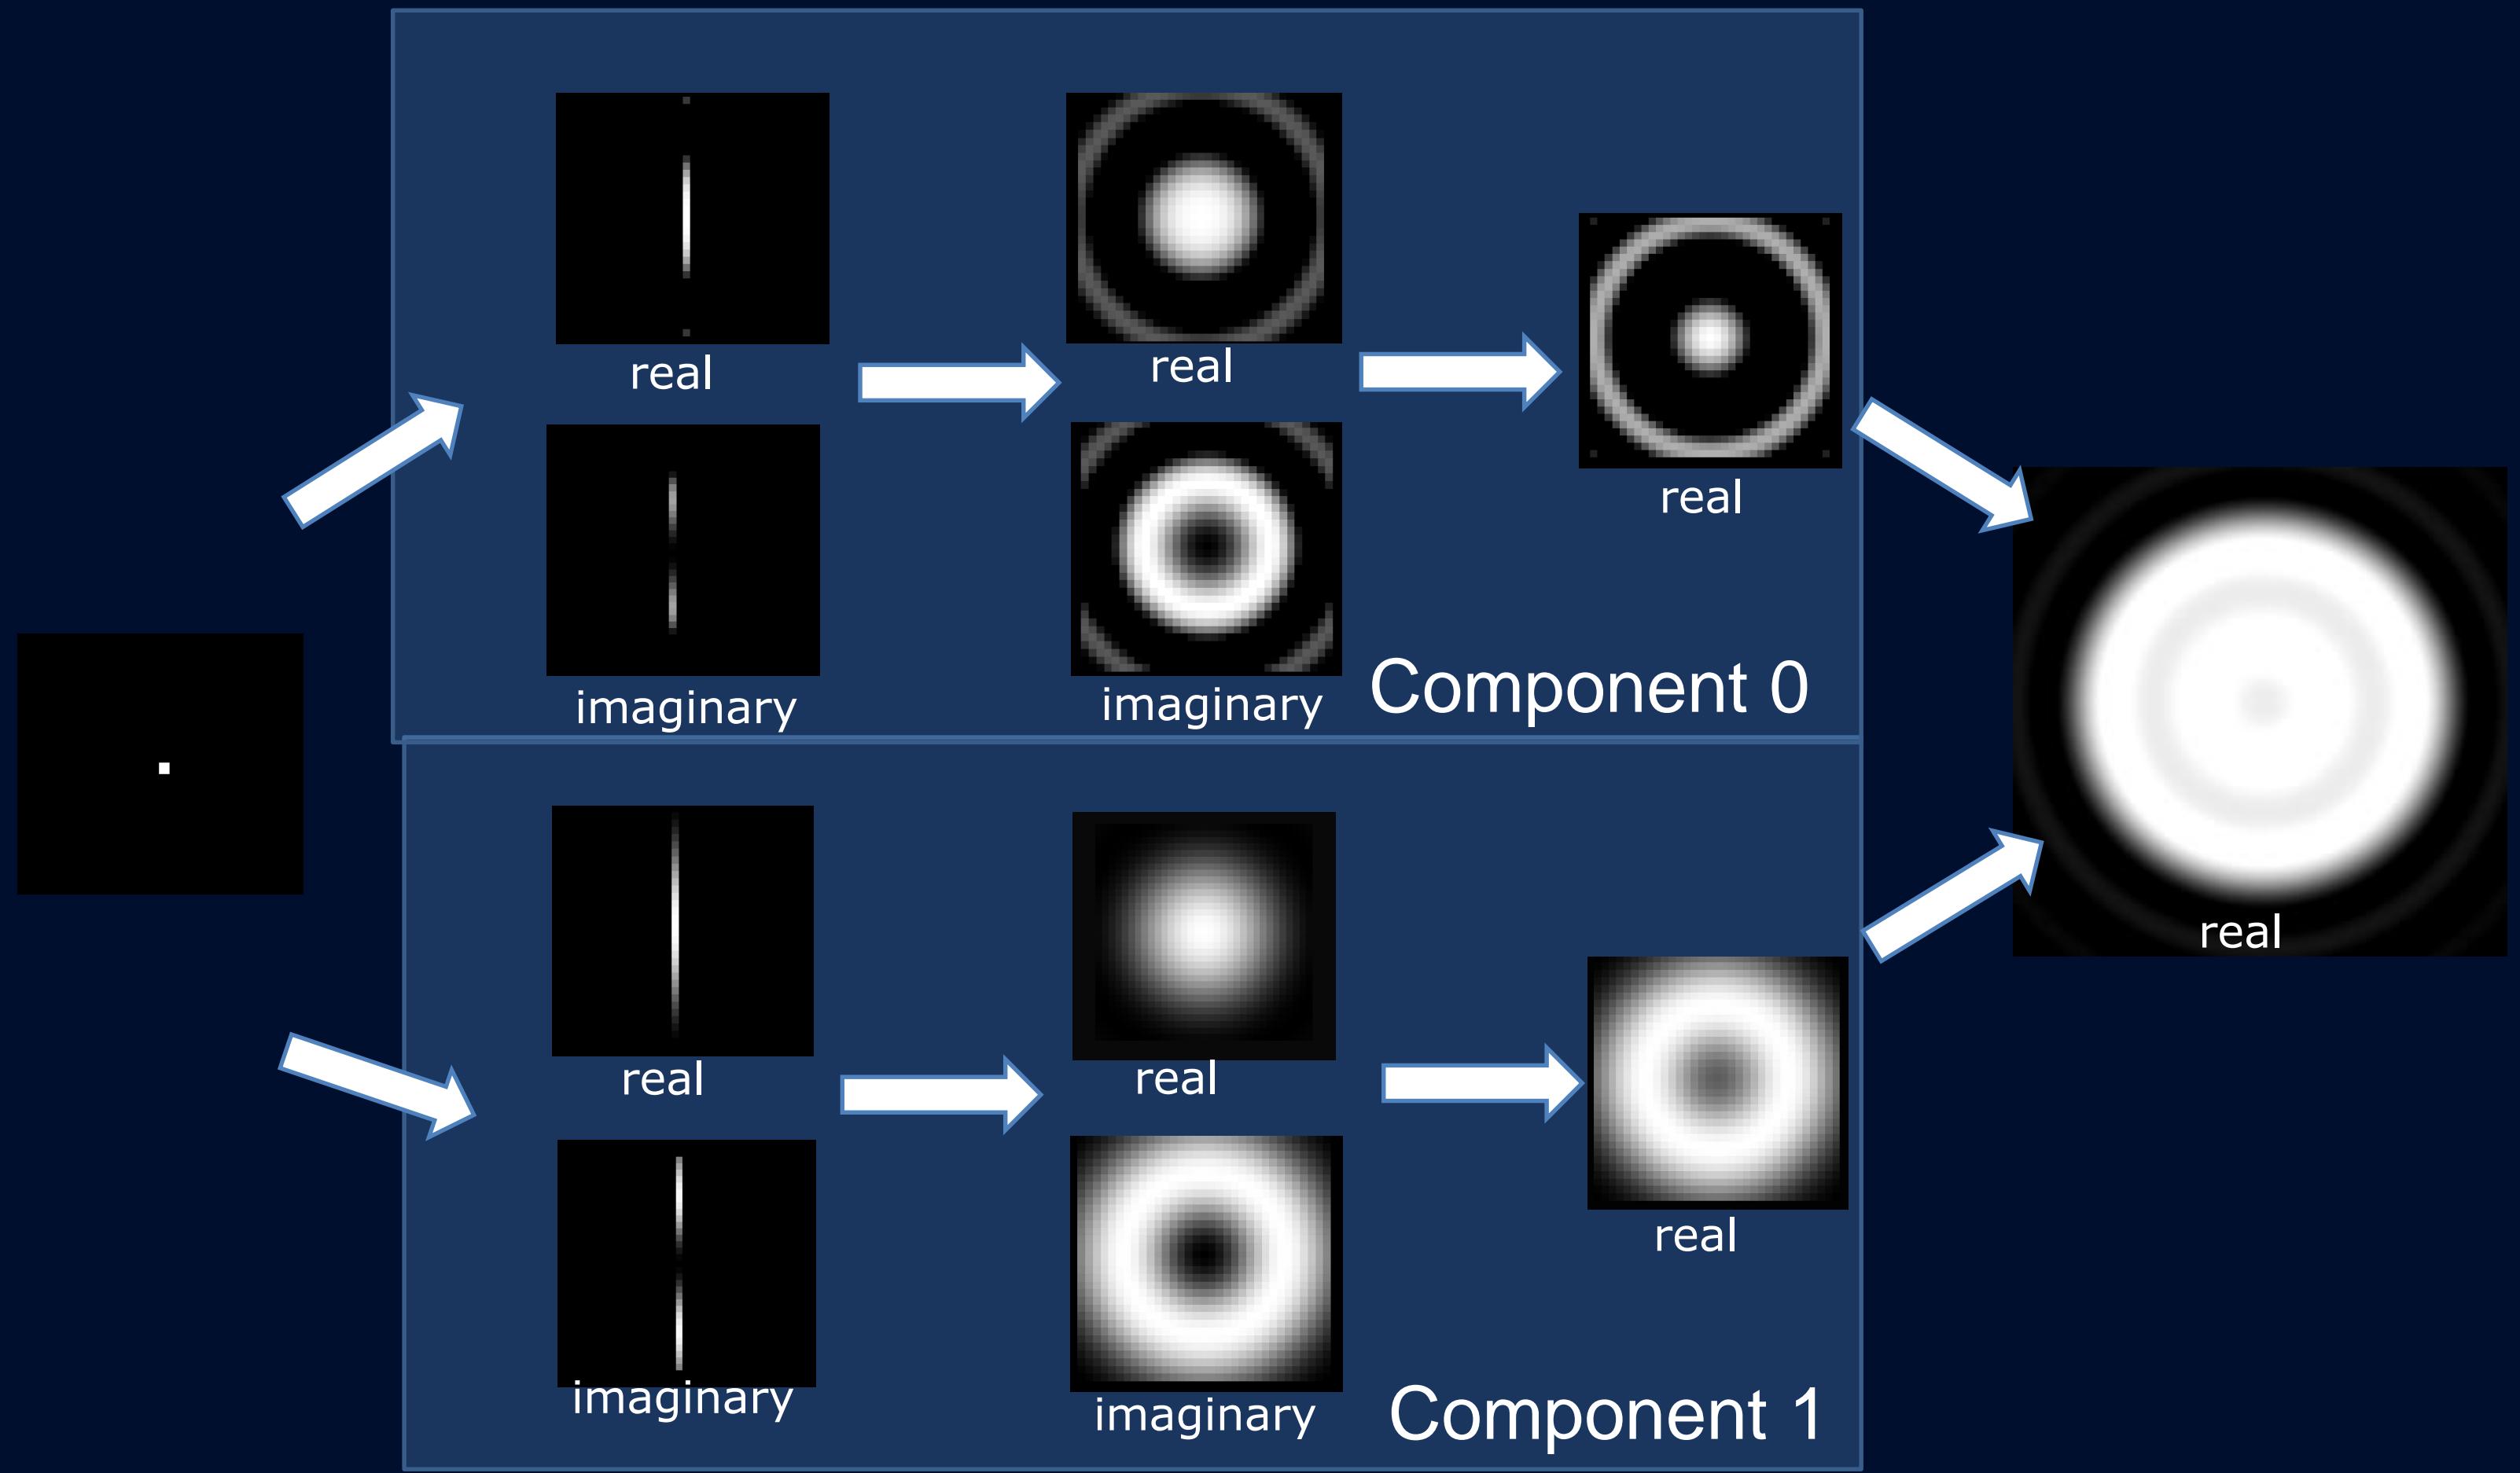
Two Component Filter

- We compute the filter the same way as before, but now with 2 components



Component	a	b	A	B
c0	-0.886528	5.268909	0.411259	-0.548794
c1	-1.960518	1.558213	0.513282	4.561110

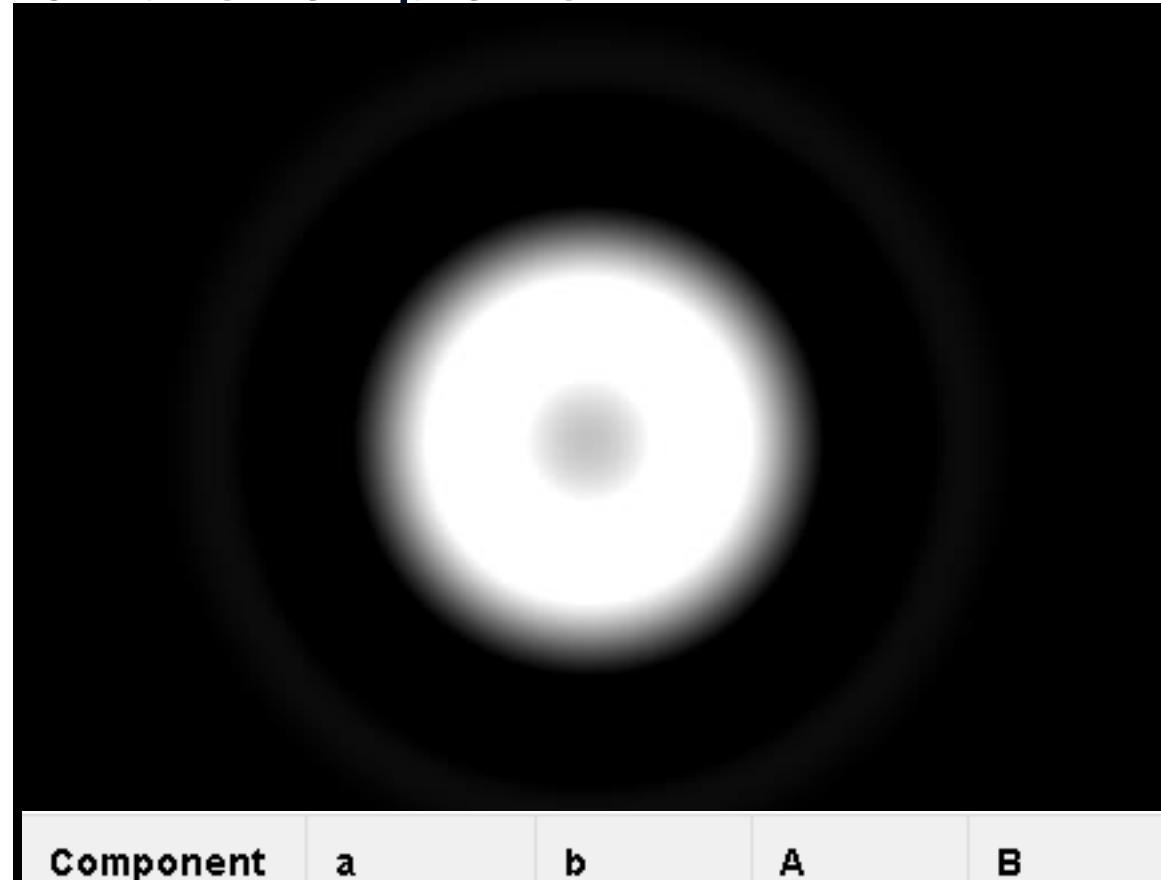




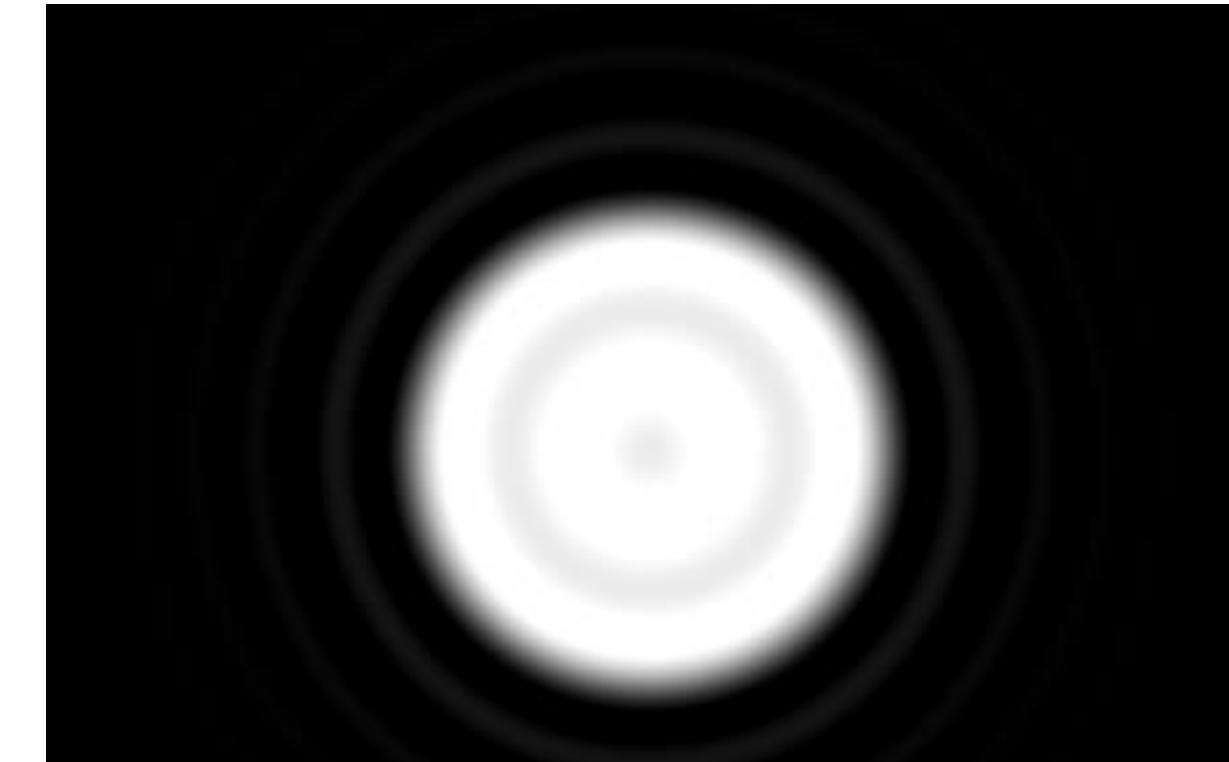
Algorithm

Circular DoF

- Low quality (1 component on left) vs High quality (2 components on the right). We use the low quality for the near blur plane, and the high quality for the far plane.



Component	a	b	A	B
c0	-0.862325	1.624835	0.767583	1.862321



Component	a	b	A	B
c0	-0.886528	5.268909	0.411259	-0.548794
c1	-1.960518	1.558213	0.513282	4.561110

$$F(x) = e^{-ax^2}(\cos(bx^2) + i \sin(bx^2))$$
$$\text{Color}(x) = A * F_{real}(x) + B * F_{imaginary}(x)$$



RGB and CoC



Downsample
PixelShader

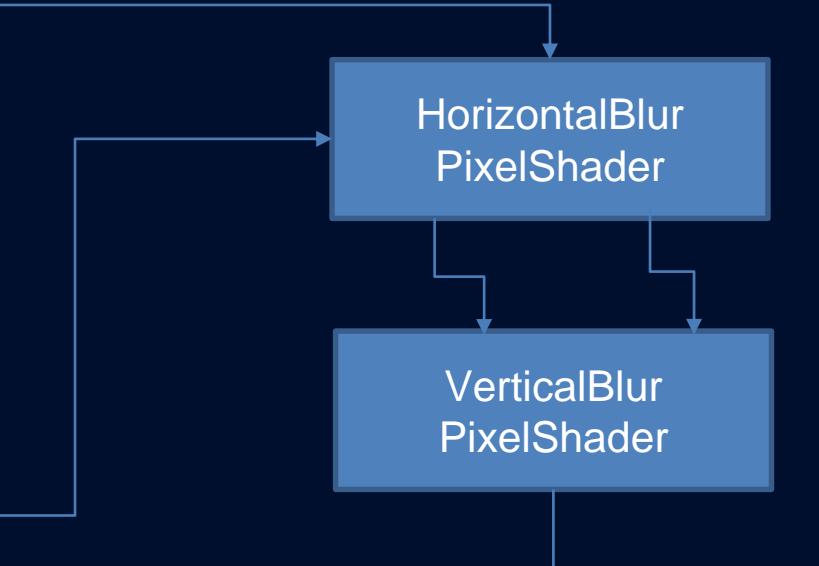
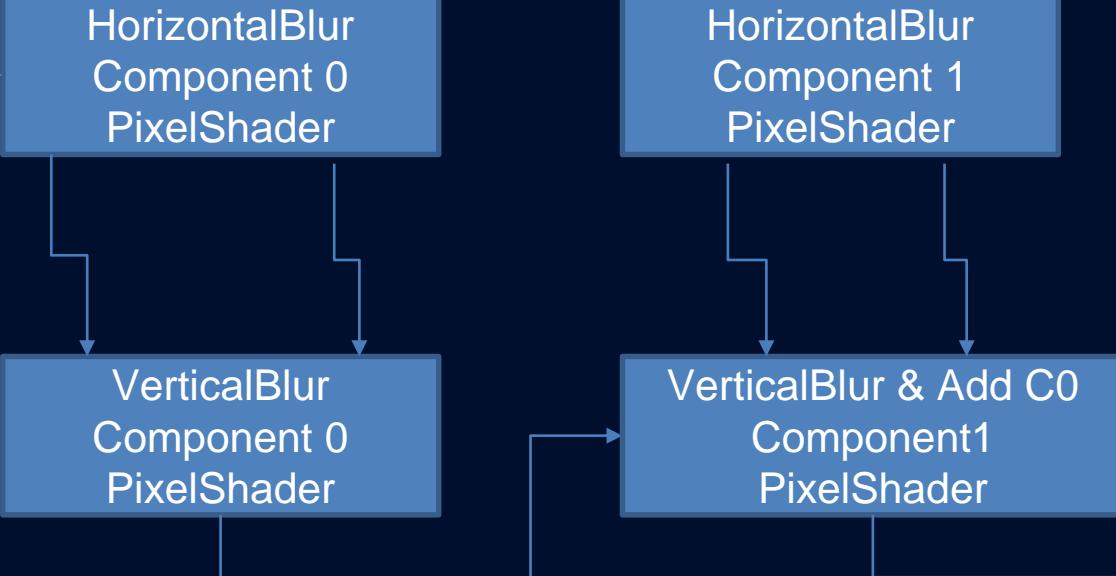
RGB



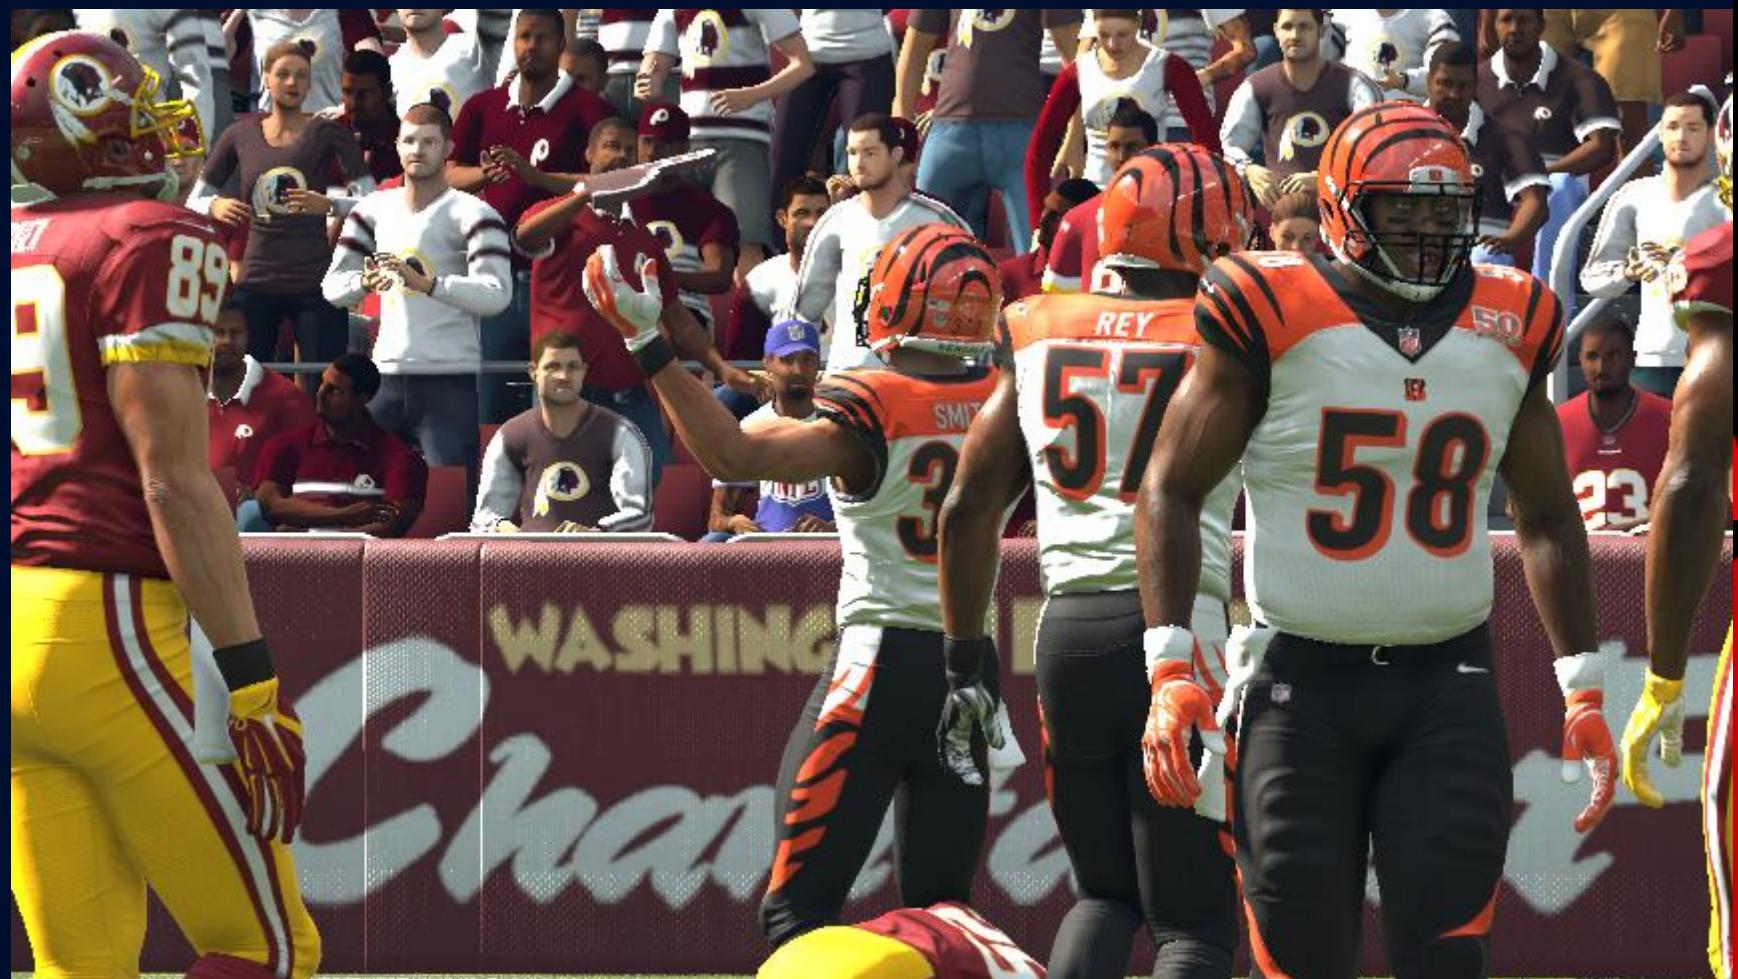
RGB



TileCocPass
ComputeShader



- Start with clear image and circle of confusion (near and far)



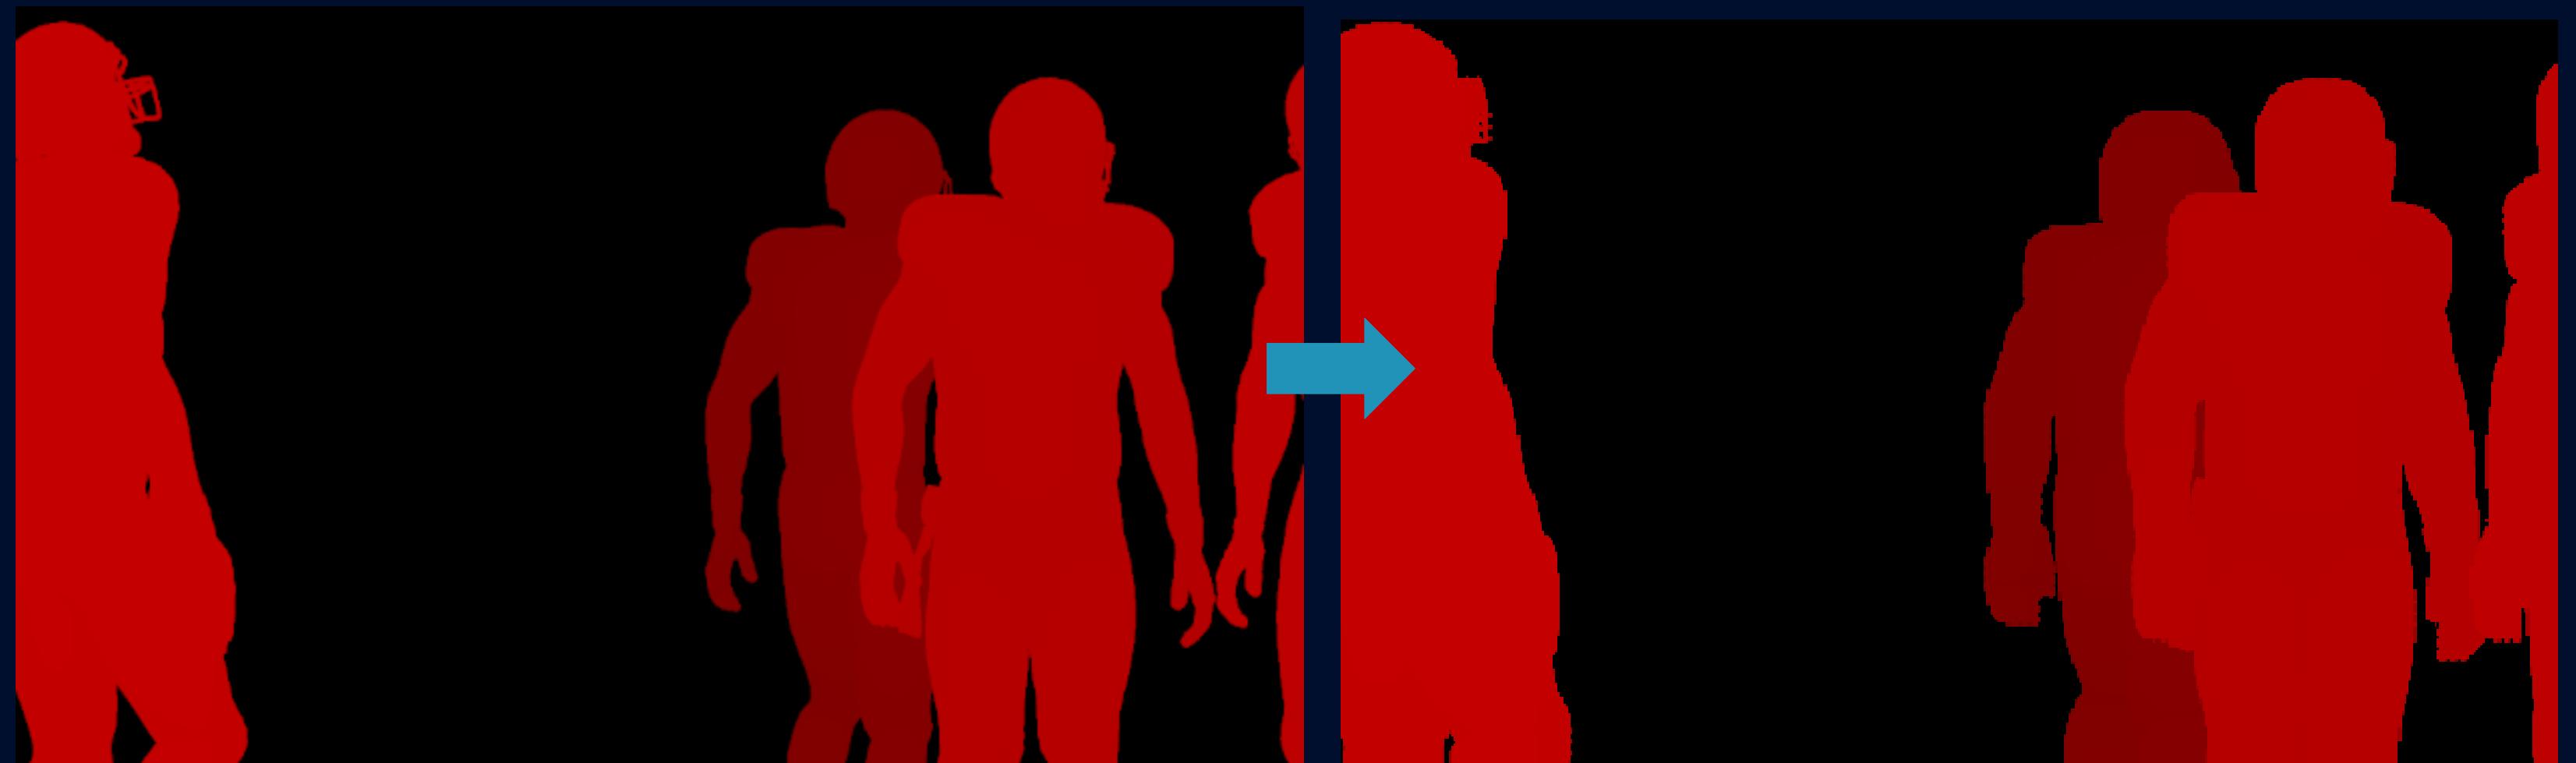
- In our case, the clear image is a 10 bit lighting buffer.
- The output would be a blurred 10 bit buffer.

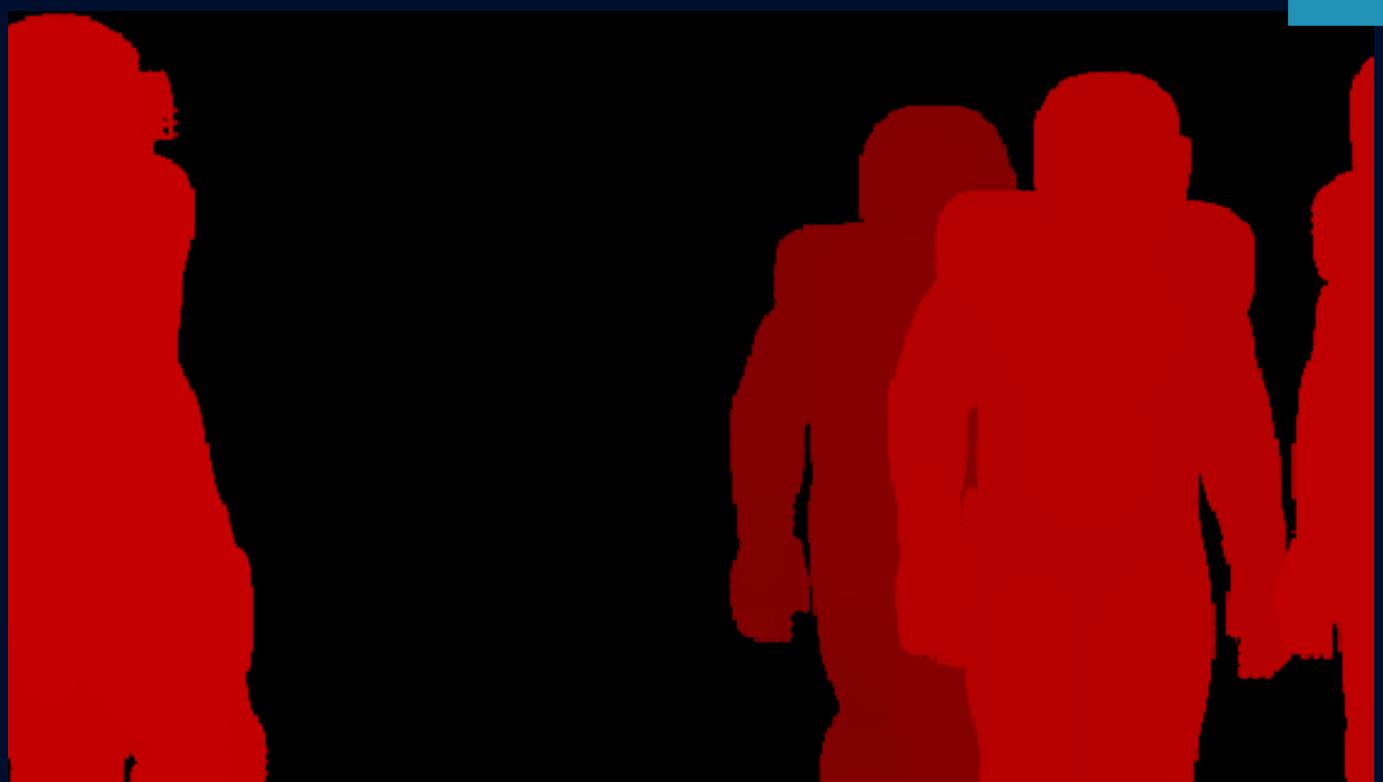
- Split main image into near and far by premultiplying circle of confusion

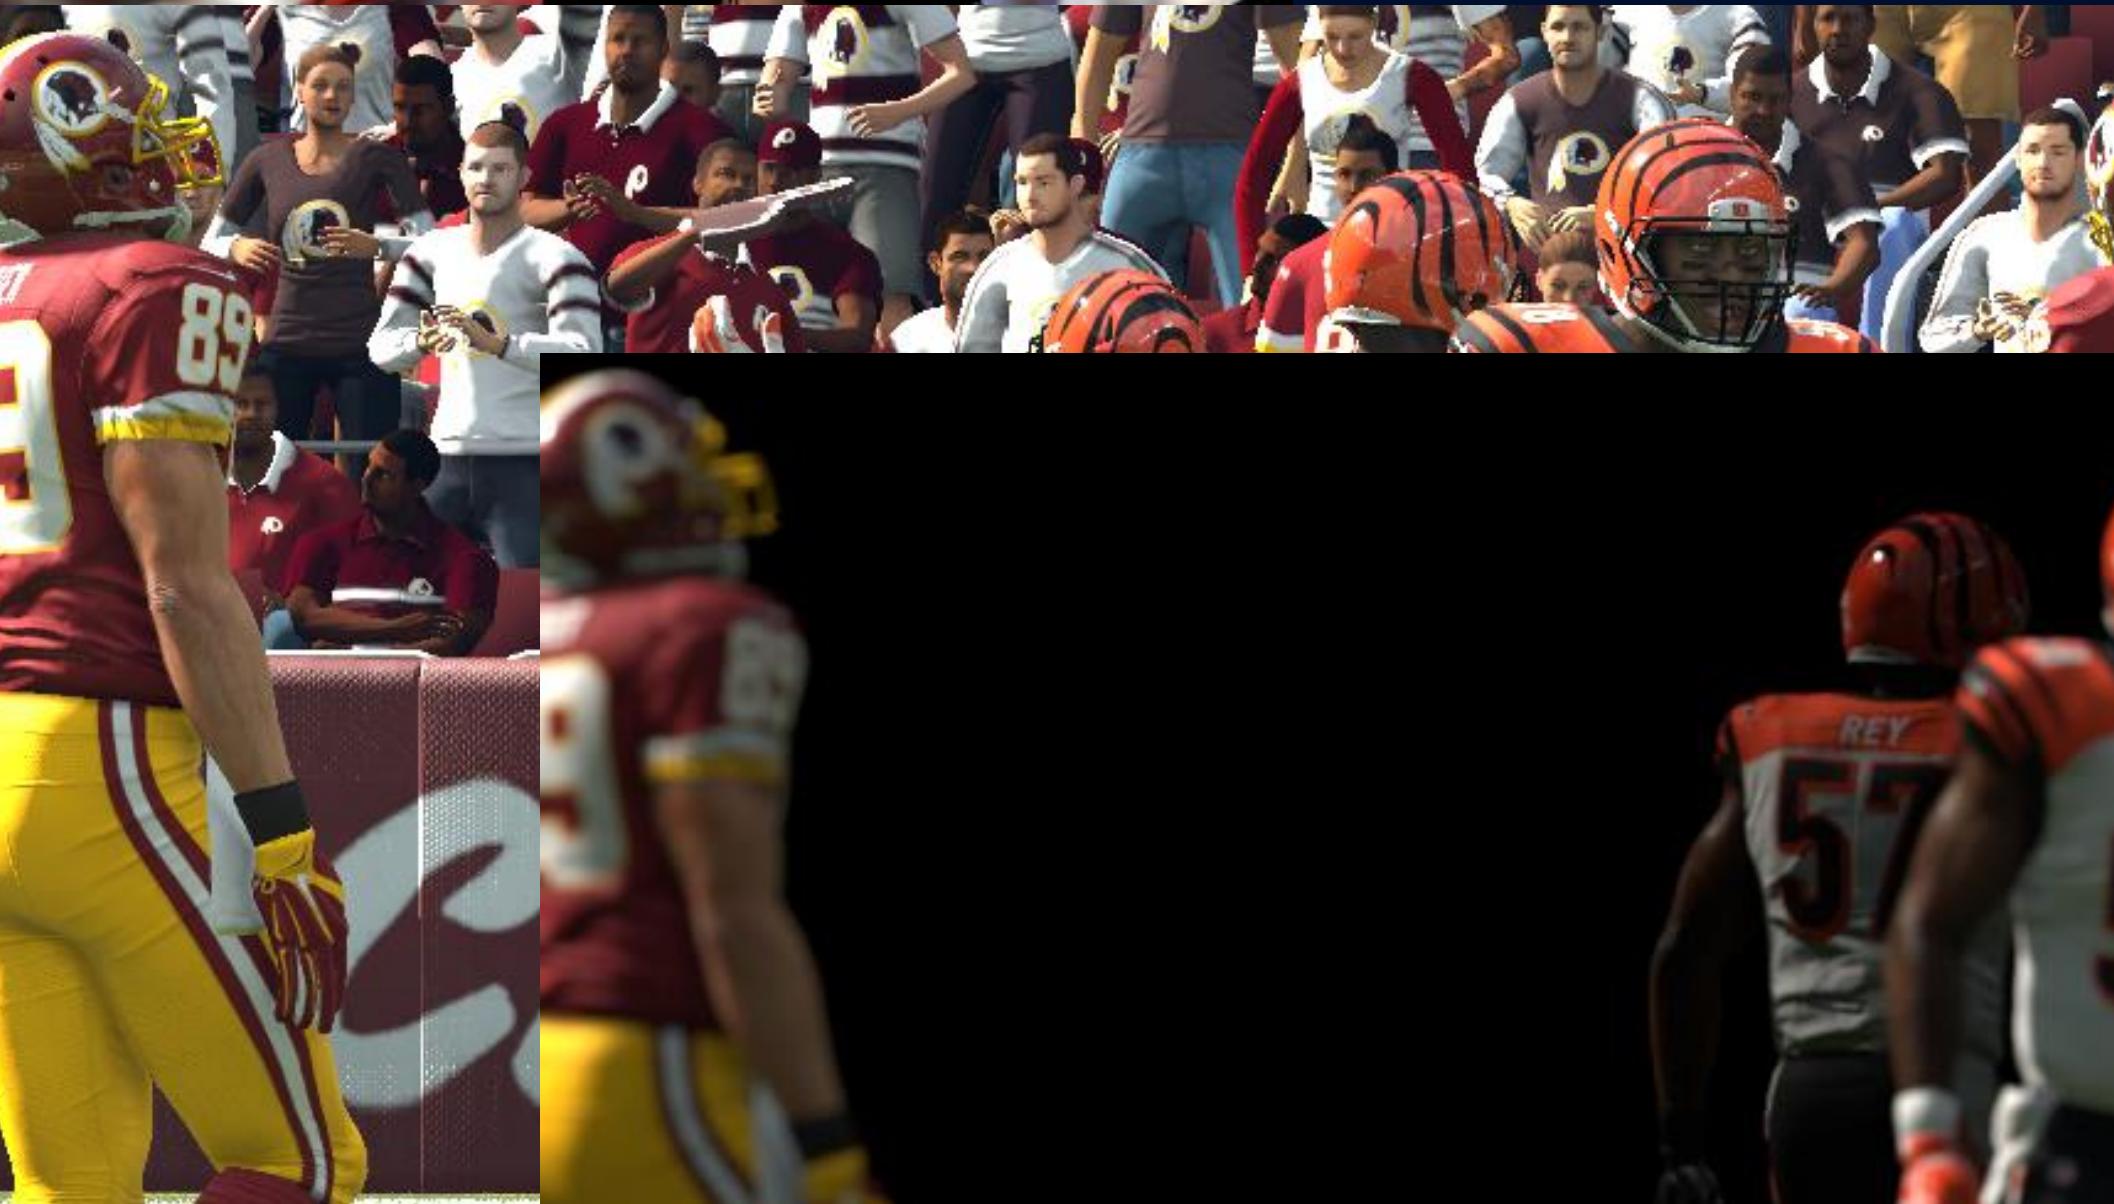




Tile the near COC (MAX within a certain tile size to get edge bleeding)







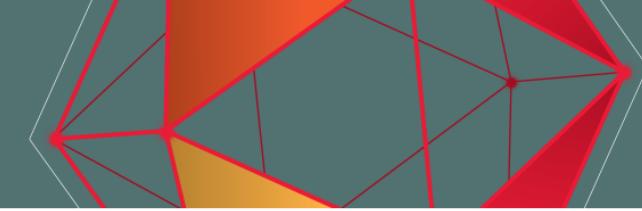




Artifacts

- Occluded circles get split in 'half' due to separable nature.
- Very subtle artifact





Artifacts

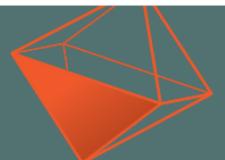
- Ghosting
- Can be reduced by biasing blending
(jumping to blur image as fast as possible)





Performance

- GPU (sprite dof 1080p, half res)
 - 9.98ms on XB1
 - 7.65ms on PS4
- GPU (1/4 res of 1080p, 2 components for far):
 - 1.7ms on XB1
 - 1.3ms on PS4
- GPU (1/2 res of 1080p, 1 component):
 - 3.4ms on XB1
 - 2.7ms on PS4





Performance (additional info)

- Limiting occupancy on xb1 and ps4 for downsampling pass of coc and color (full res to quarter res)
- Downsampling pass is massively vmem bound, ends up trashing the cache.
- Solution? Limit the occupancy!, can make it run as fast!
- XB1:
 - `#define __XBOX_LIMIT_OCCUPANCY_WITH_HARD_LIMIT 1`
- PS4:
 - `#pragma argument(minvgprcount=84)`
 - `#pragma argument(targetoccupancy=3)`
 - `#pragma argument(scheduler=latency)`
- special thanks to [Tomasz Stachowiak \[@h3r2tic \]](#)





Additional Perf Opportunities

- Explore amortization – essentially less samples are required for smaller CoC radii. We can precompute a set of filter weights, for different radius ranges, and pick them dynamically.
 - Would not improve performance on a fully blurred image.
 - Would improve performance for areas fully clear of the image.
- Combine near and far
 - Essentially have only one shader for horizontal and vertical passes
 - Use MRTs to output different values of near and far
 - Might have to explore manual occupancy hints to preserve vmem cache coherency





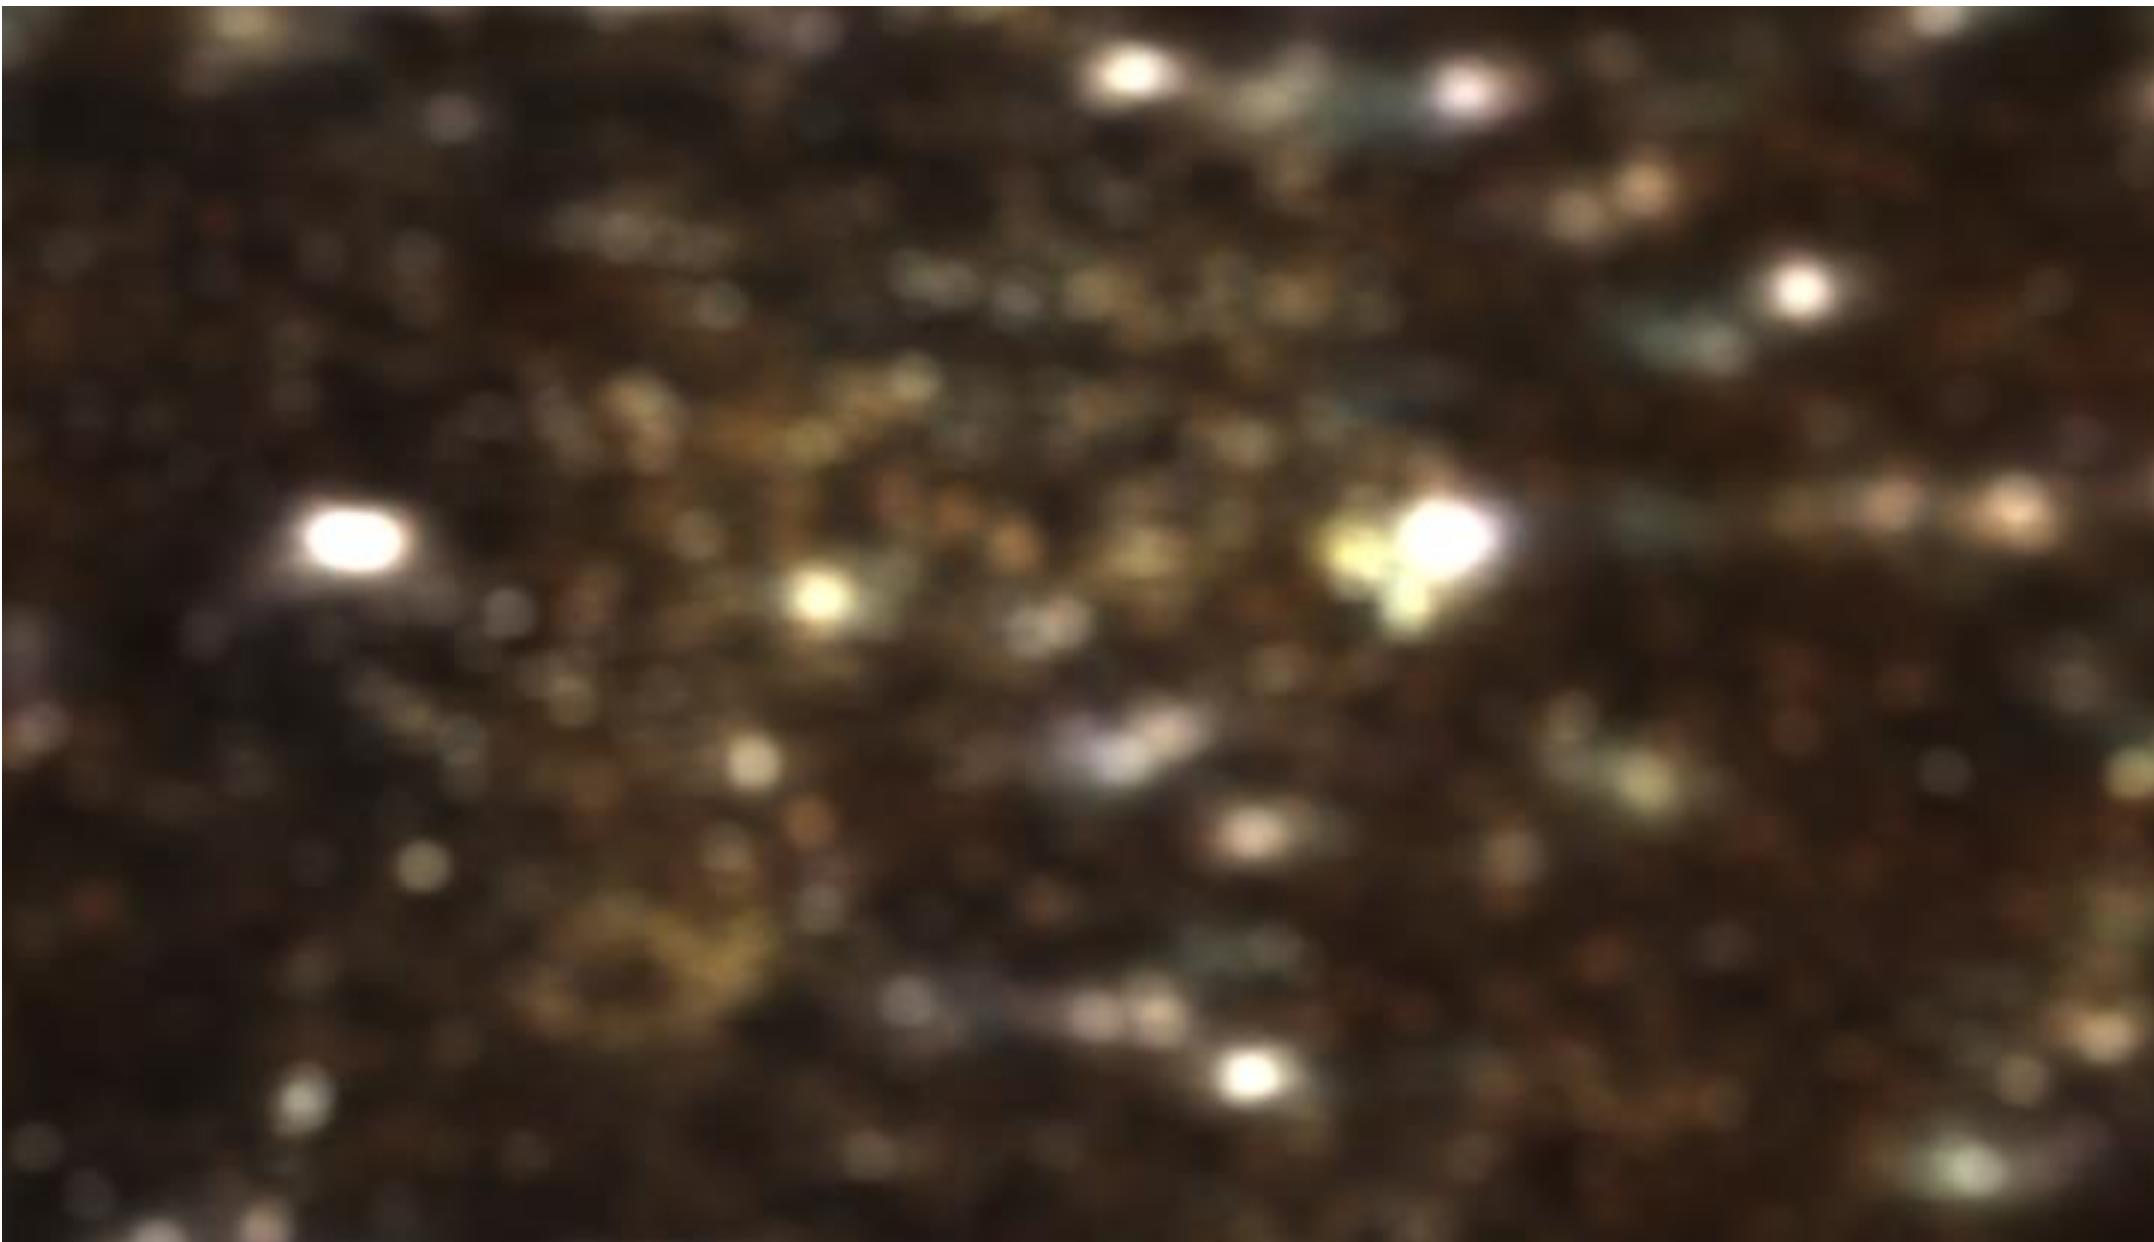
Additional Visual Improvements

- Transparency: instead of using transparent depth to shift COC, use multiple render planes / buckets and composite these.
- More on Ghosting
 - improved performance gains and do the pass in full resolution (see previous slide!)
 - dynamically compute pixel to ratio bias, and use scene information such as pixel luminance to automatically 'jump' to the next blur plane.



Shader Toy Example

<https://www.shadertoy.com/view/Xd2BWc>



PreFiltering

- Used a filter generator algorithm to precompute the filter
- Madden uses a 68 pixels (in $\frac{1}{4}$ res $r = 8$) diameter filter!
- It uses 2 component for far blur and 1 component for near blur.
- <https://github.com/kecho/CircularDofFilterGenerator>
 - A lite python version of the filter generator can be found here



Sources

- CSC algorithm blog post. (Olli Niemitalo, 2010) <http://yehar.com/blog/?p=1495>
- Five Rendering ideas from BF3 and NFS: e run, (Electronic Arts, Siggraph 2011)
<http://www.slideshare.net/DICEStudio/five-rendering-ideas-from-battlefield-3-need-for-speed-the-run>





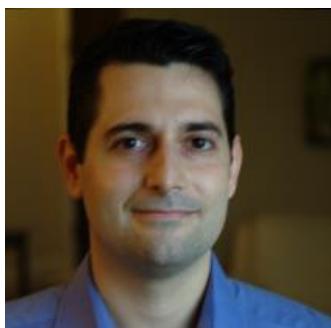
Credits



Kleber Garcia - Render Engineer, Frostbite



Karolyn Boulanger – Render Engineer, EA Sports



Arthur Homs – Principal Engineer, Microsoft

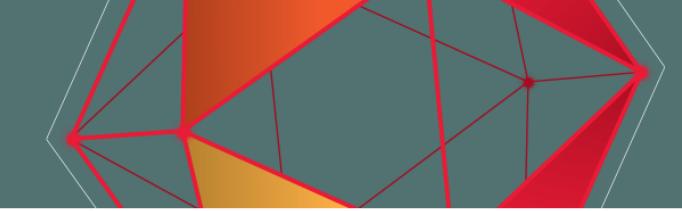


Ollie Niemitalo – Mathematician, Signal processing scientist.

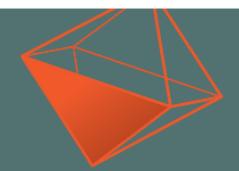


Q & A





Appendix – Mathematical derivations.

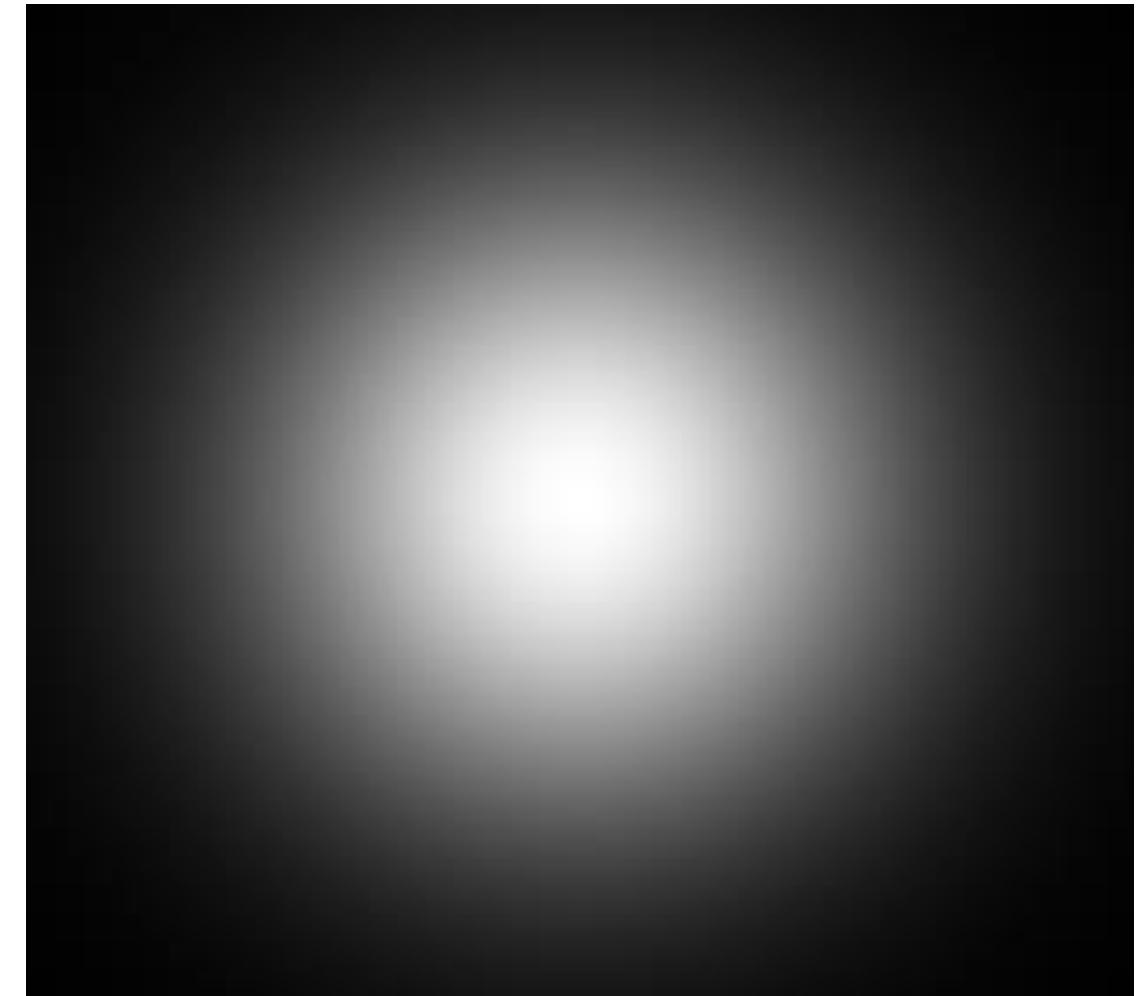


$F(x)$ filter derivation

- A separable filter $F(x)$, is separable when:

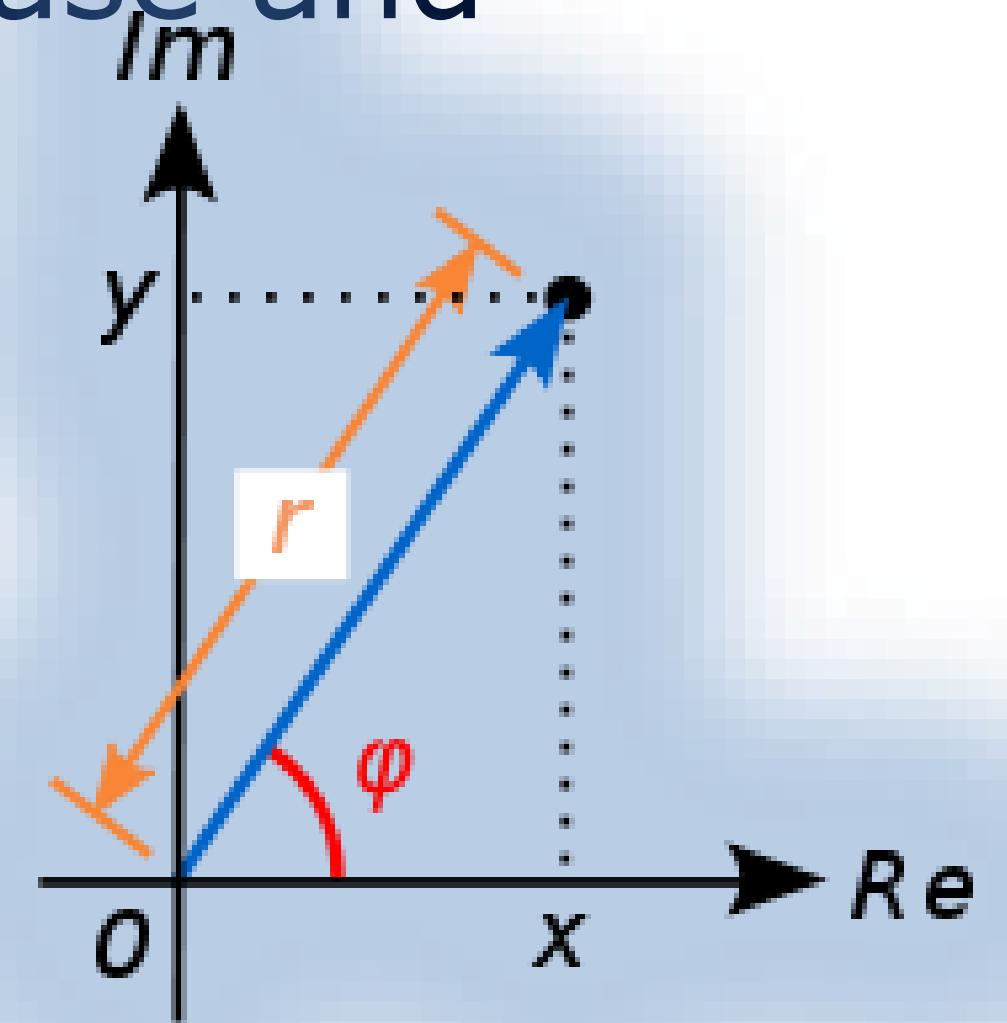
- $F(\sqrt{x^2+y^2}) = F(x) * F(y)$

$F(x) = e^{-x^2}$ Gaussian function has this property! Therefore that's why is separable



$F(x)$ filter derivation

- An imaginary number has a phase and envelope:
- Imaginary number can be written as $x + iy$
- Or: $r(\cos\varphi + i \sin\varphi)$
- Or: $re^{i\varphi}$



$F(x)$ filter derivation

- Let $F(x)$ be a complex function.
- Let $|F(x)|$ be the magnitude (r in the previous slide)
- Let $\arg(F(x))$ be the envelope (angle φ)
- $F(x)$ can be written as:
 - $F(x) = |F(x)| * (\cos(\arg(F(x))) + i * (\sin(\arg(F(x)))))$

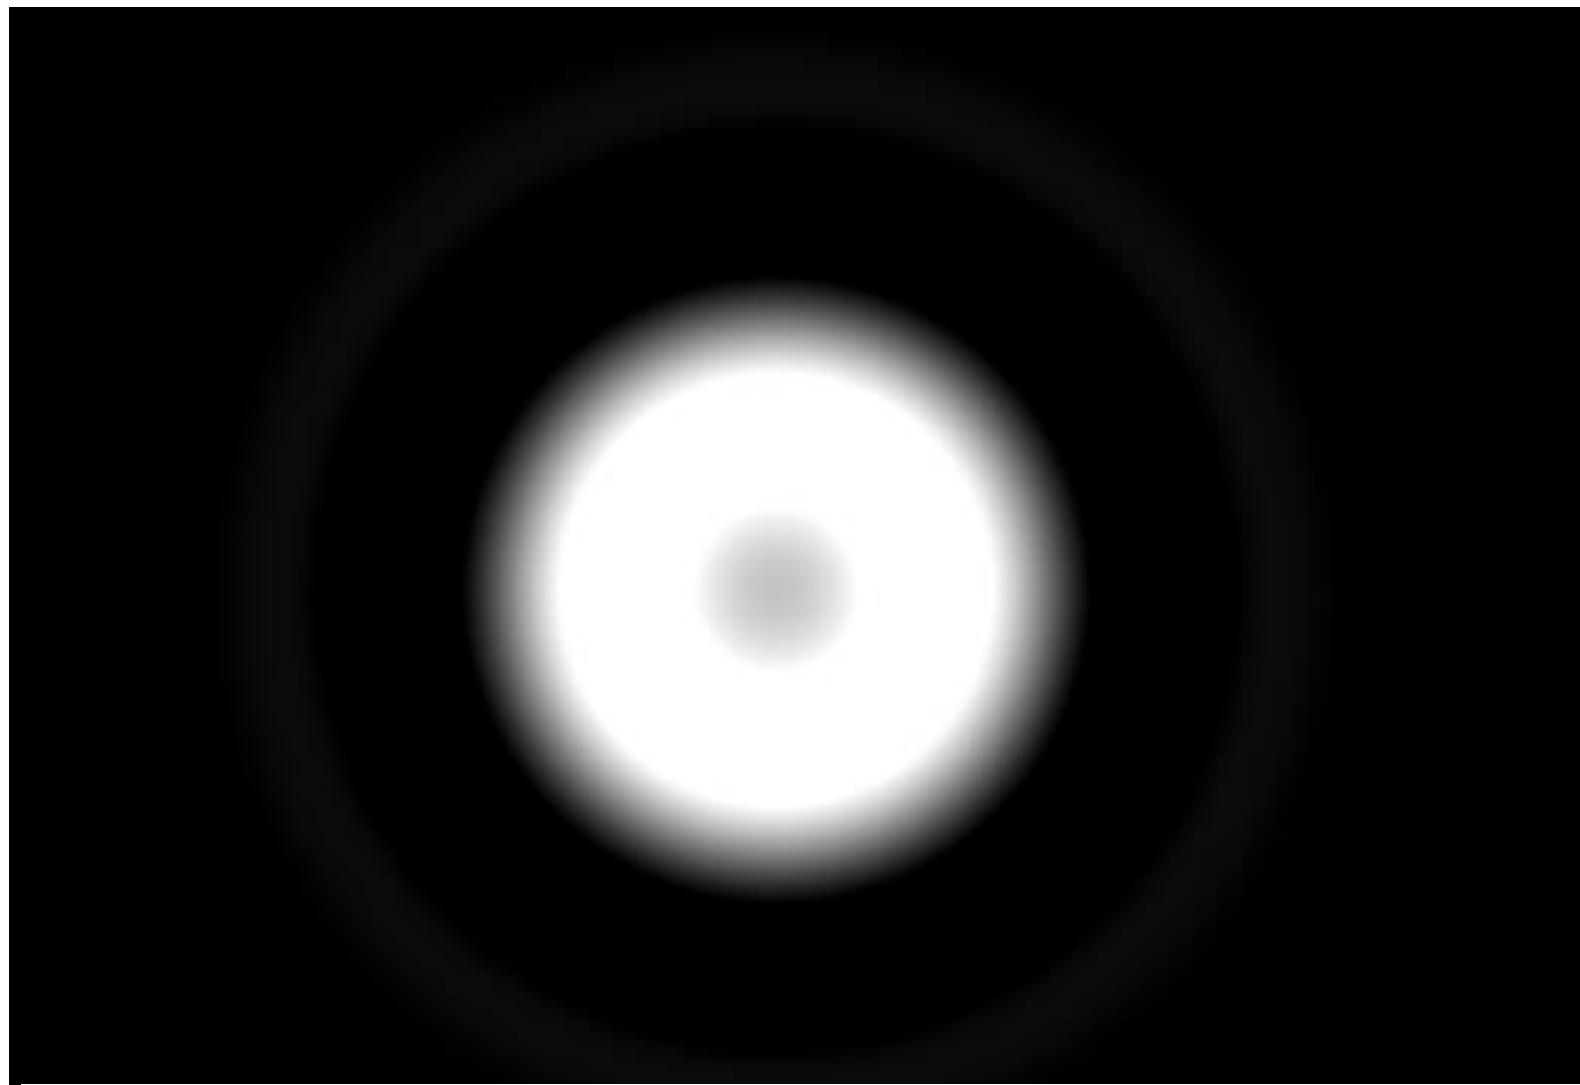
$F(x)$ filter derivation

- $F(x) = |F(x)| * (\cos(\arg(F(x))) + i * (\sin(\arg(F(x))))$
- Assume $F(x)$ is separable.
- Hence $|F(x)|$ must be a Gaussian function
 - $|F(x)| = e^{-a*x^2}$
 - $\arg(F(\sqrt{x^2 + y^2})) = \arg(F(x)) + \arg(F(y))$
 - therefore $\arg(F(x)) = bx^2$

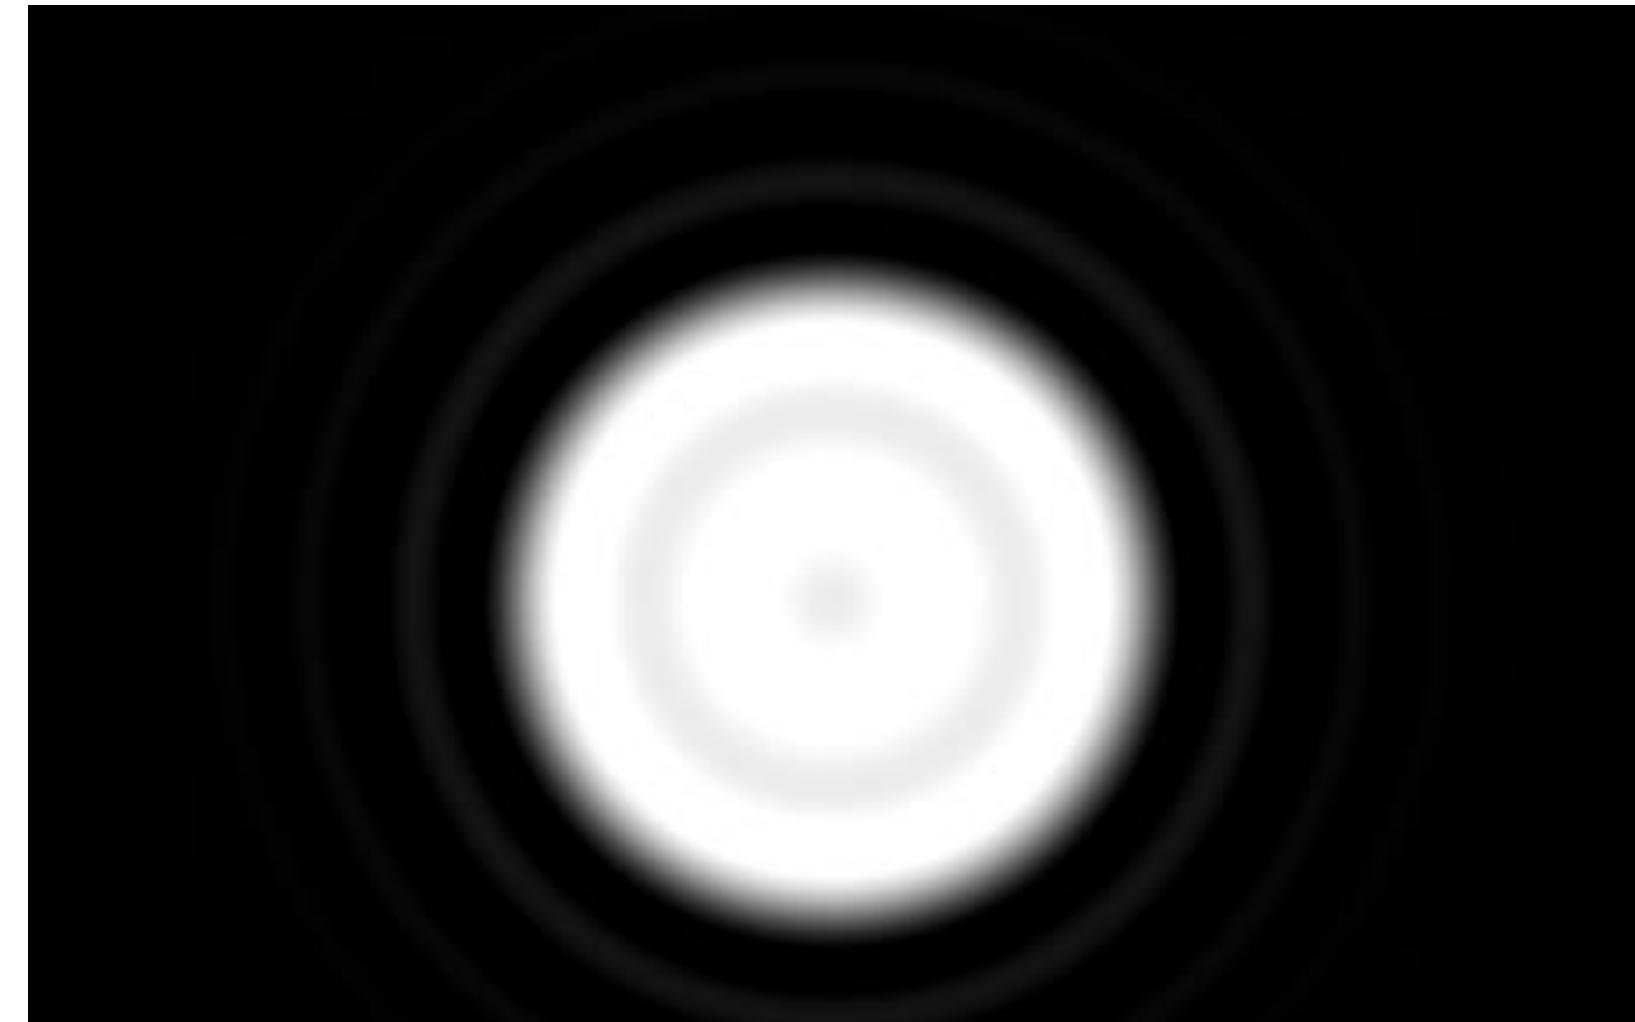
$F(x)$ filter derivation

- Replacing the previous terms, we get
 - $F(x) = e^{-ax^2} * (\cos(b*x^2) + i * \sin(b*x^2))$
- Arbitrary circles can be achieved a weighted sum of imaginary and real elements.
- Final filter kernel function becomes:
 - $F_{final}(x) = A * F_{real}(x) + B * F_{imaginary}(x)$
- A final sum of these components will give us a convoluted color.

$F(x)$ filter derivation



Component	a	b	A	B
C0	-0.862325	1.624835	0.767583	1.862321



Component	a	b	A	B
C0	-0.886528	5.268909	0.411259	-0.548794
C1	-1.960518	1.558213	0.513282	4.561110

Bracketing the filter

- How can we maximize bit precision? Bracketing and squeezing the filter to produce numbers in the $[0,1]$ domain.

Bracketing the filter

- Let $0 < x < N$, where N is the max pixel width.
- Assume we have an arbitrary G kernel with the following properties:
 - $\sum_{x=1}^N G(x) = V$
 - $O = \text{Min}(G(x))$
 - $S = \sum_{x=1}^N \{G(x) - O\}$
- We can then transform the kernel G into the bracketed kernel G' , which can be defined as:
 - $G'(x) = \frac{G(x) - O}{S}$
- We can then store coefficients O and S for $G'(x)$

Bracketing the filter

- Let I be a 1 dimensional (for simplicity) image, 16 bit rgba buffer for our final storage.
- Let I' be a temp storage, which can only store numbers from 0 to 1 (10 bit rgba buffer)
- Let J be our initial image.
- Lets now try to convolve J using $G'(x)$ and store it in $I'[x]$
 - Since we know O , and S we can store
 - $I'[w] = \sum_{x=1}^N J[w]G'(x)$ <- Lets instead store the bracketed version, and separately keep track of the kernel values O and S .
- $I'[i]$ is not quite what we planned though! We want to take $I'[i]$ and convert it to the equivalent of $I[w] = \sum_{x=1}^N I[w]G(x)$

Bracketing the filter

- Now, we know that I' contains our bracketed filter values. When we sample from I' , we can convert to the actual non bracketed by applying some inverse operations.
- We know $I[w] = \sum_{x=1}^N J[w]G(x)$
- Here is how we convert I' into I :
 - We know the definition of I' . We can expand I' algebraically into
 - $I'[w] = \sum_{x=1}^N J[w] * \left[\frac{G[i]-O}{S} \right]$
 - Means we can do some algebra and define I as
 - $I[w] = \sum_{x=1}^N \left(J[w] * \left[\frac{G[i]-O}{S} \right] * S \right) + \sum_{x=1}^N J[w]O$
 - $I[w] = \sum_{x=1}^N (I'[w] * S) + \sum_{x=1}^N J[w] * O$
 - Means that if we store O , S and Sum of all $J[w]$ s (in a separate target) we can compress the render targets into 10 bits with unbounded information.